

IESTI01 - TinyML

The Building Blocks of Deep
Learning – Part C

Prof. Marcelo Rovai

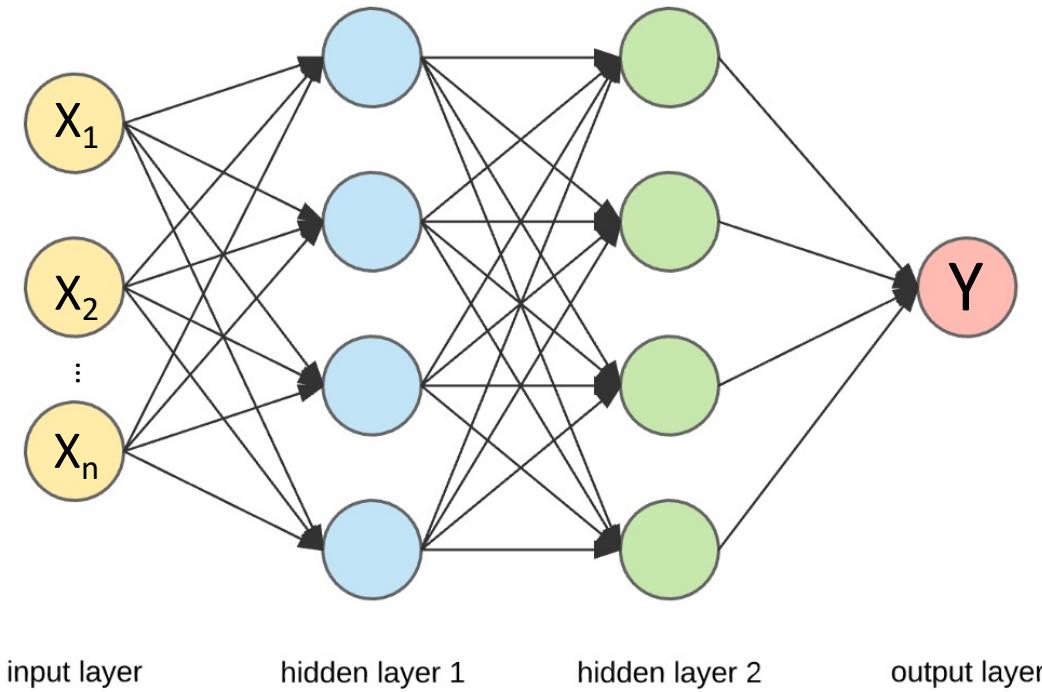
May 26th, 2021

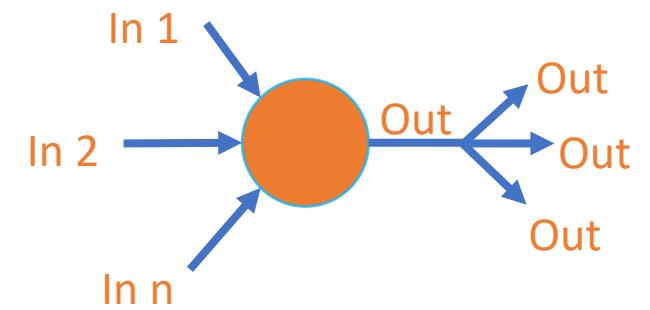
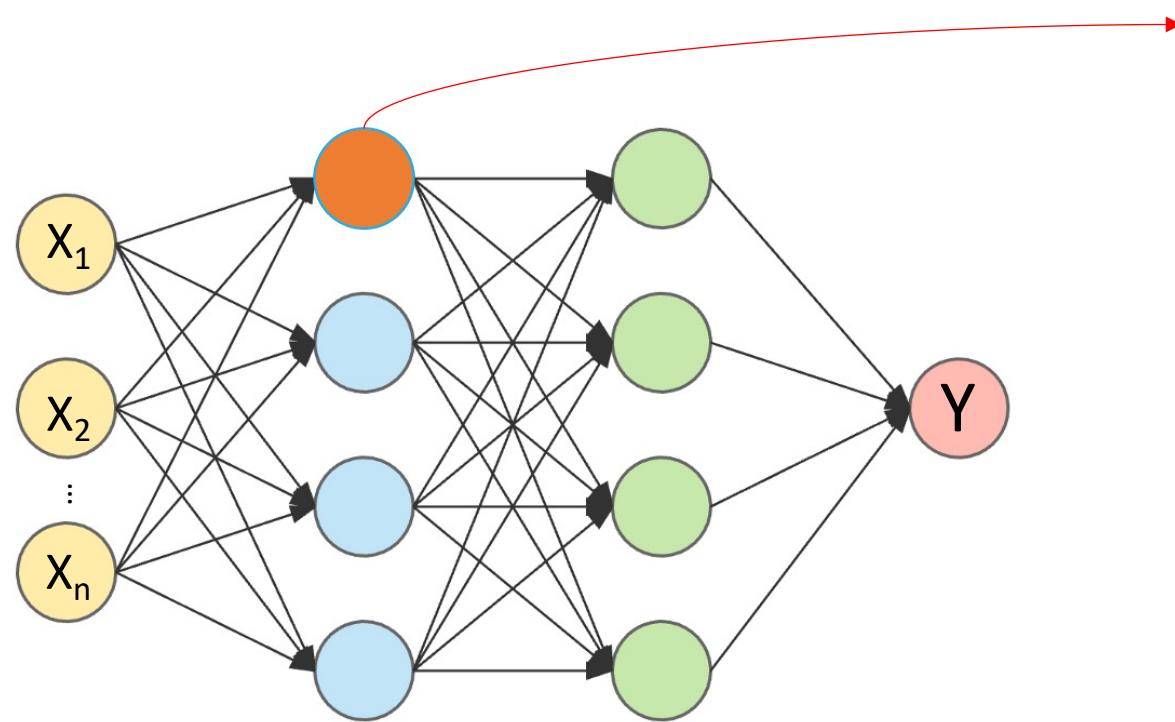


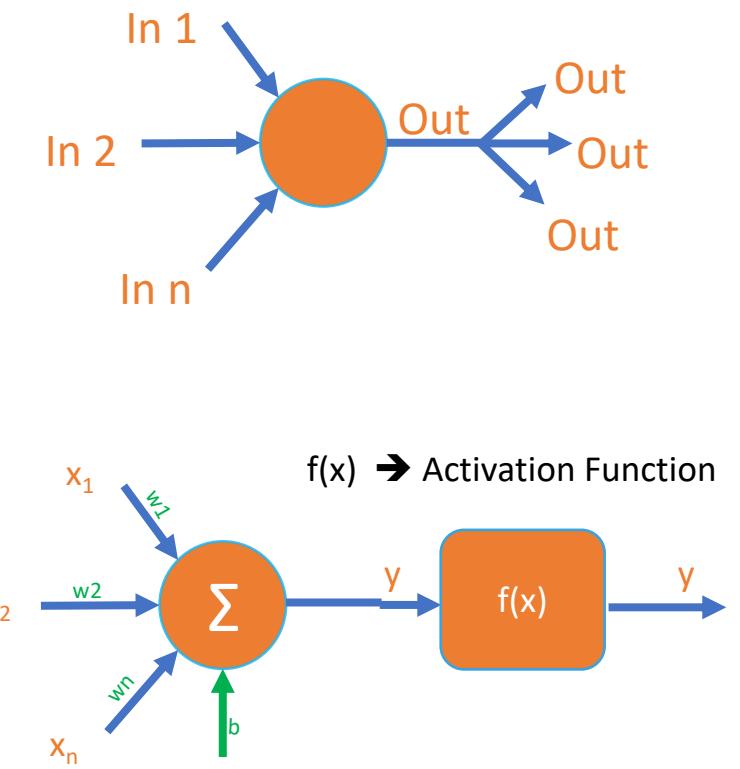
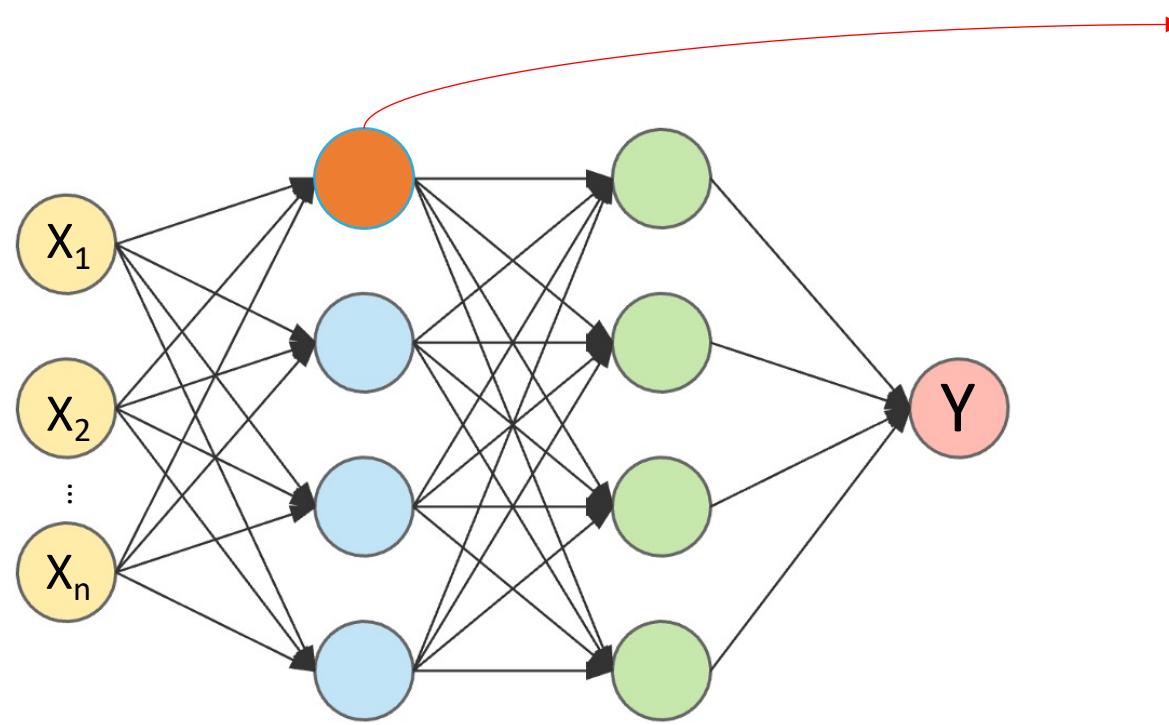
DNN Dense Neural Network

Recap

Supervised Machine Learning with DNN

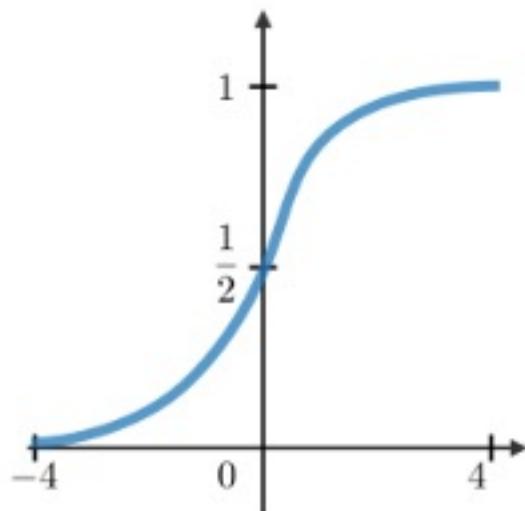
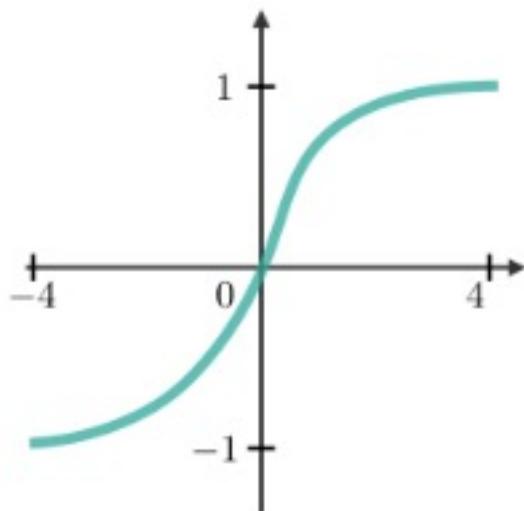
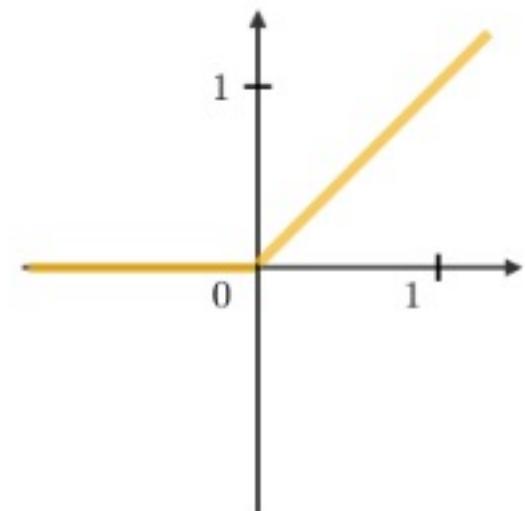


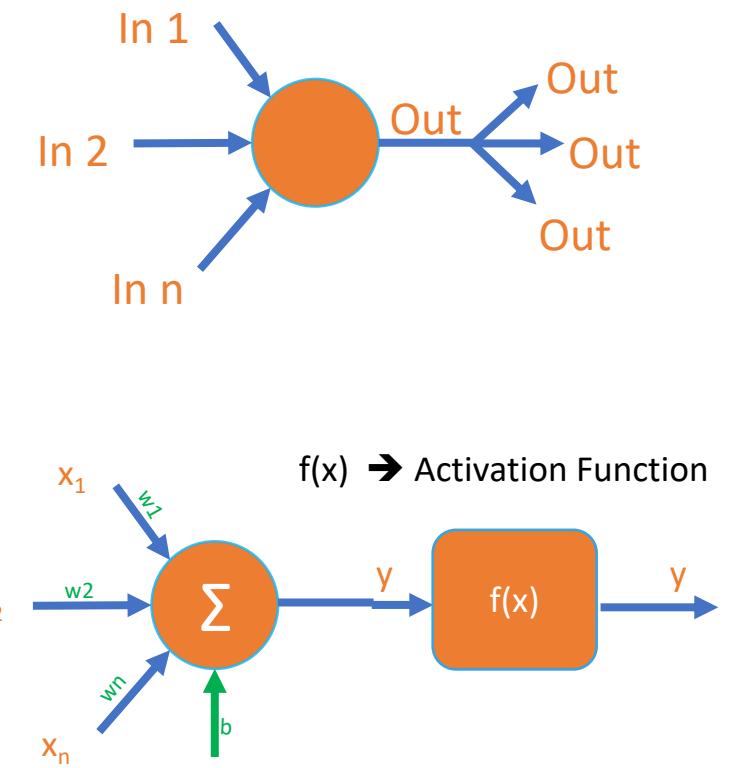
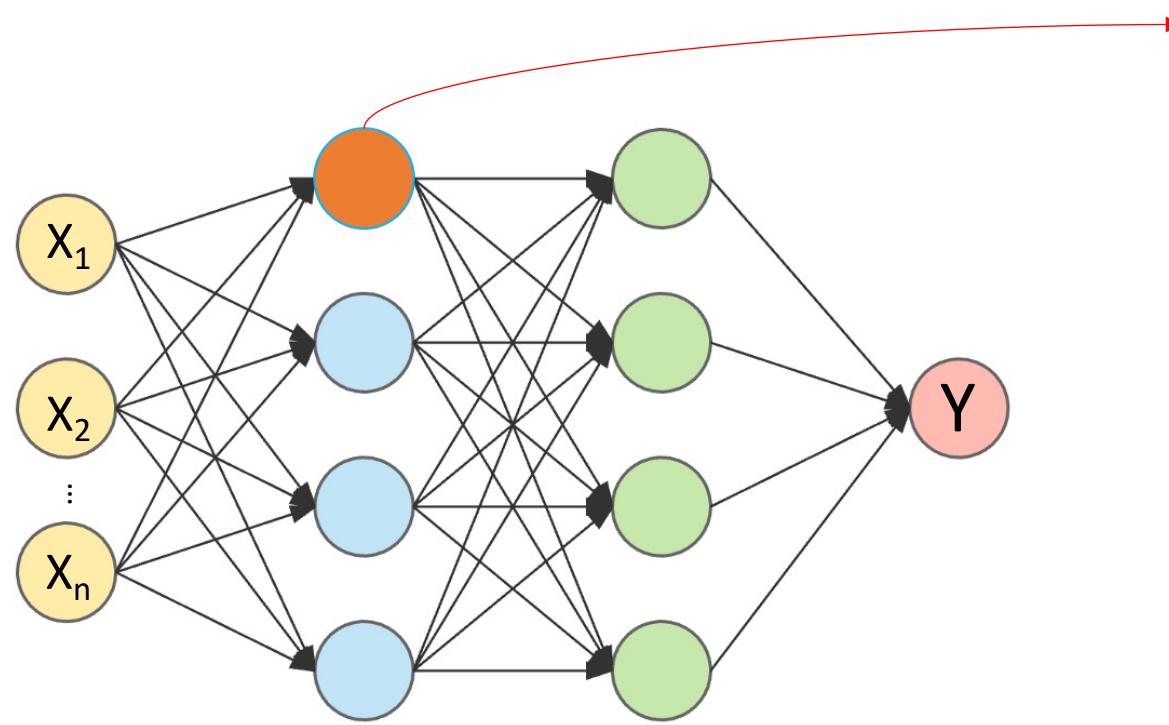




$$y = f\left(\sum_{i=1}^n x_i w_i + b\right)$$

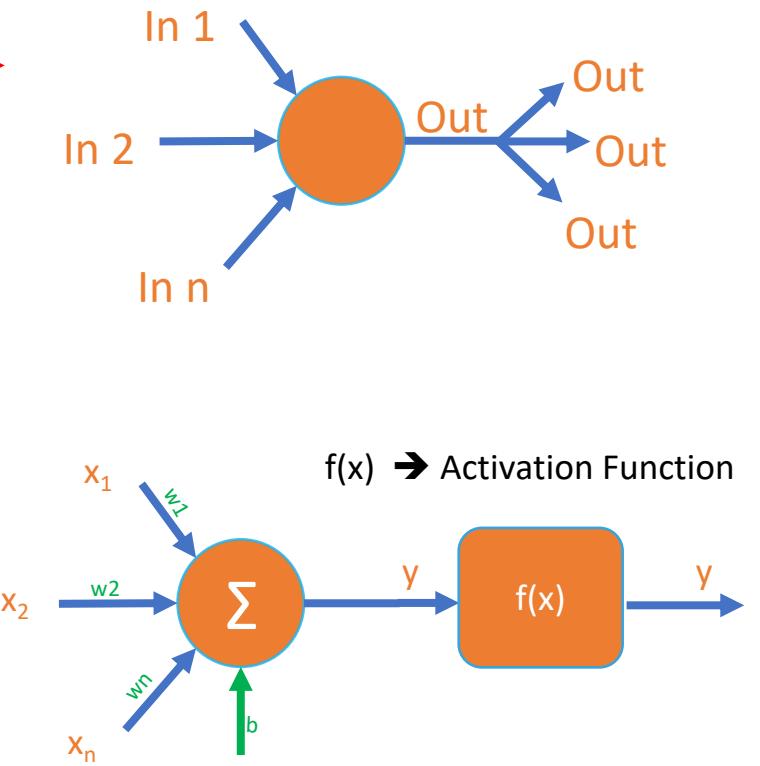
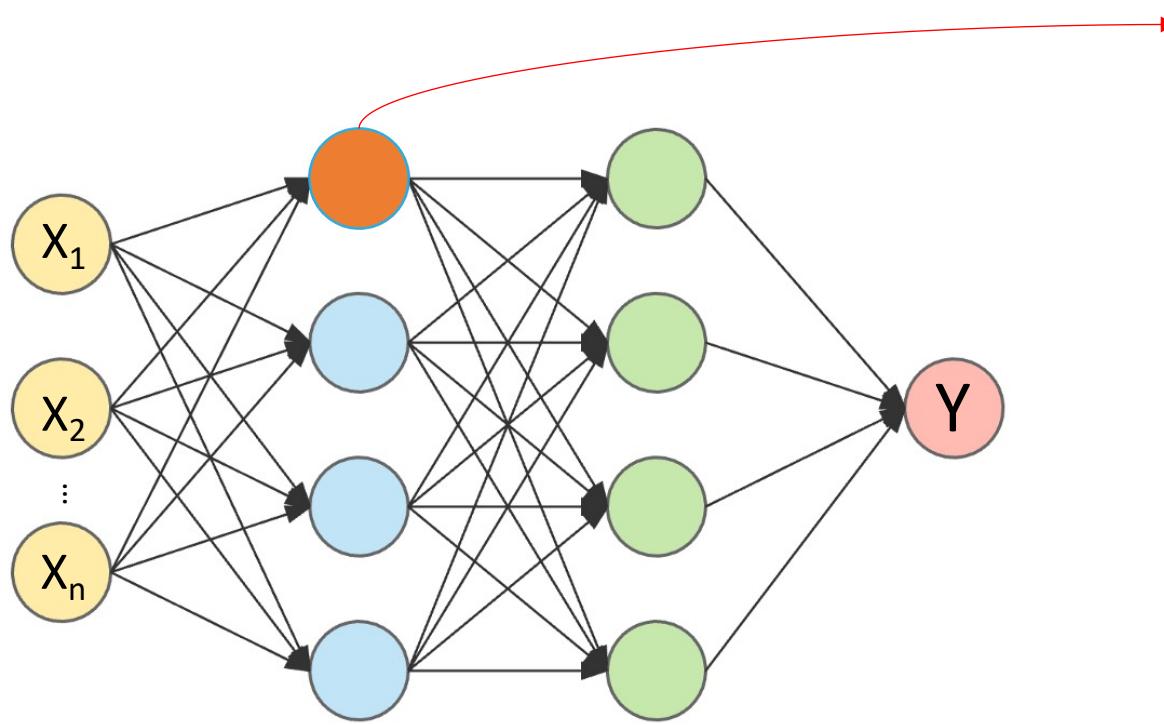
Activation Functions

Sigmoid	Tanh	RELU
$g(z) = \frac{1}{1 + e^{-z}}$	$g(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$	$g(z) = \max(0, z)$
 The graph shows the sigmoid function $y = \frac{1}{1 + e^{-x}}$. The x-axis ranges from -4 to 4, and the y-axis ranges from 0 to 1. The curve passes through the point (0, 0.5), which is marked with a vertical line. The curve approaches y=0 as x goes to negative infinity and y=1 as x goes to positive infinity.	 The graph shows the hyperbolic tangent function $y = \tanh(x)$. The x-axis ranges from -4 to 4, and the y-axis ranges from -1 to 1. The curve passes through the origin (0, 0). It is symmetric about the origin and approaches y=-1 as x goes to negative infinity and y=1 as x goes to positive infinity.	 The graph shows the Rectified Linear Unit (ReLU) function $y = \max(0, x)$. The x-axis ranges from -1 to 1, and the y-axis ranges from 0 to 1. The function is zero for all negative x values and increases linearly with a slope of 1 for all positive x values.



Parameters to be found during training
To reach minimum error

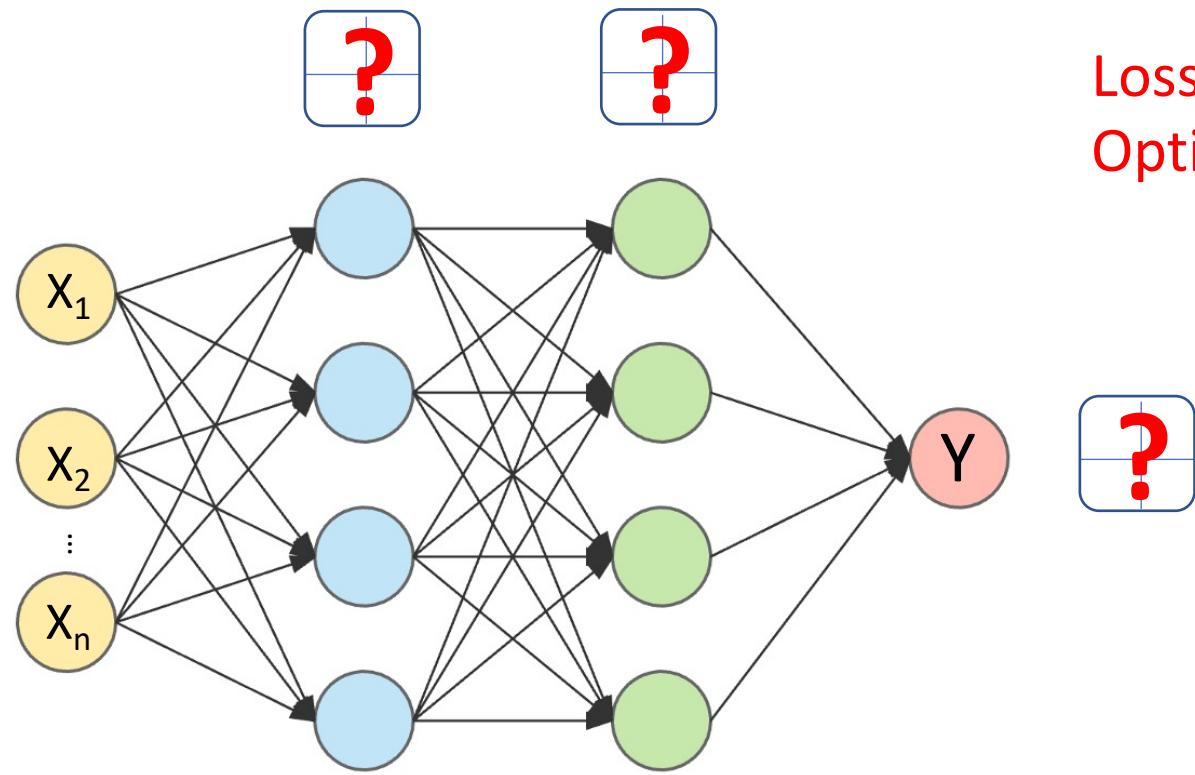
$$y = f\left(\sum_{i=1}^n x_i w_i + b\right)$$



- Error Measurement (Loss)
- Optimization

Parameters to be found during training
To reach minimum error

$$y = f\left(\sum_{i=1}^n x_i w_i + b\right)$$

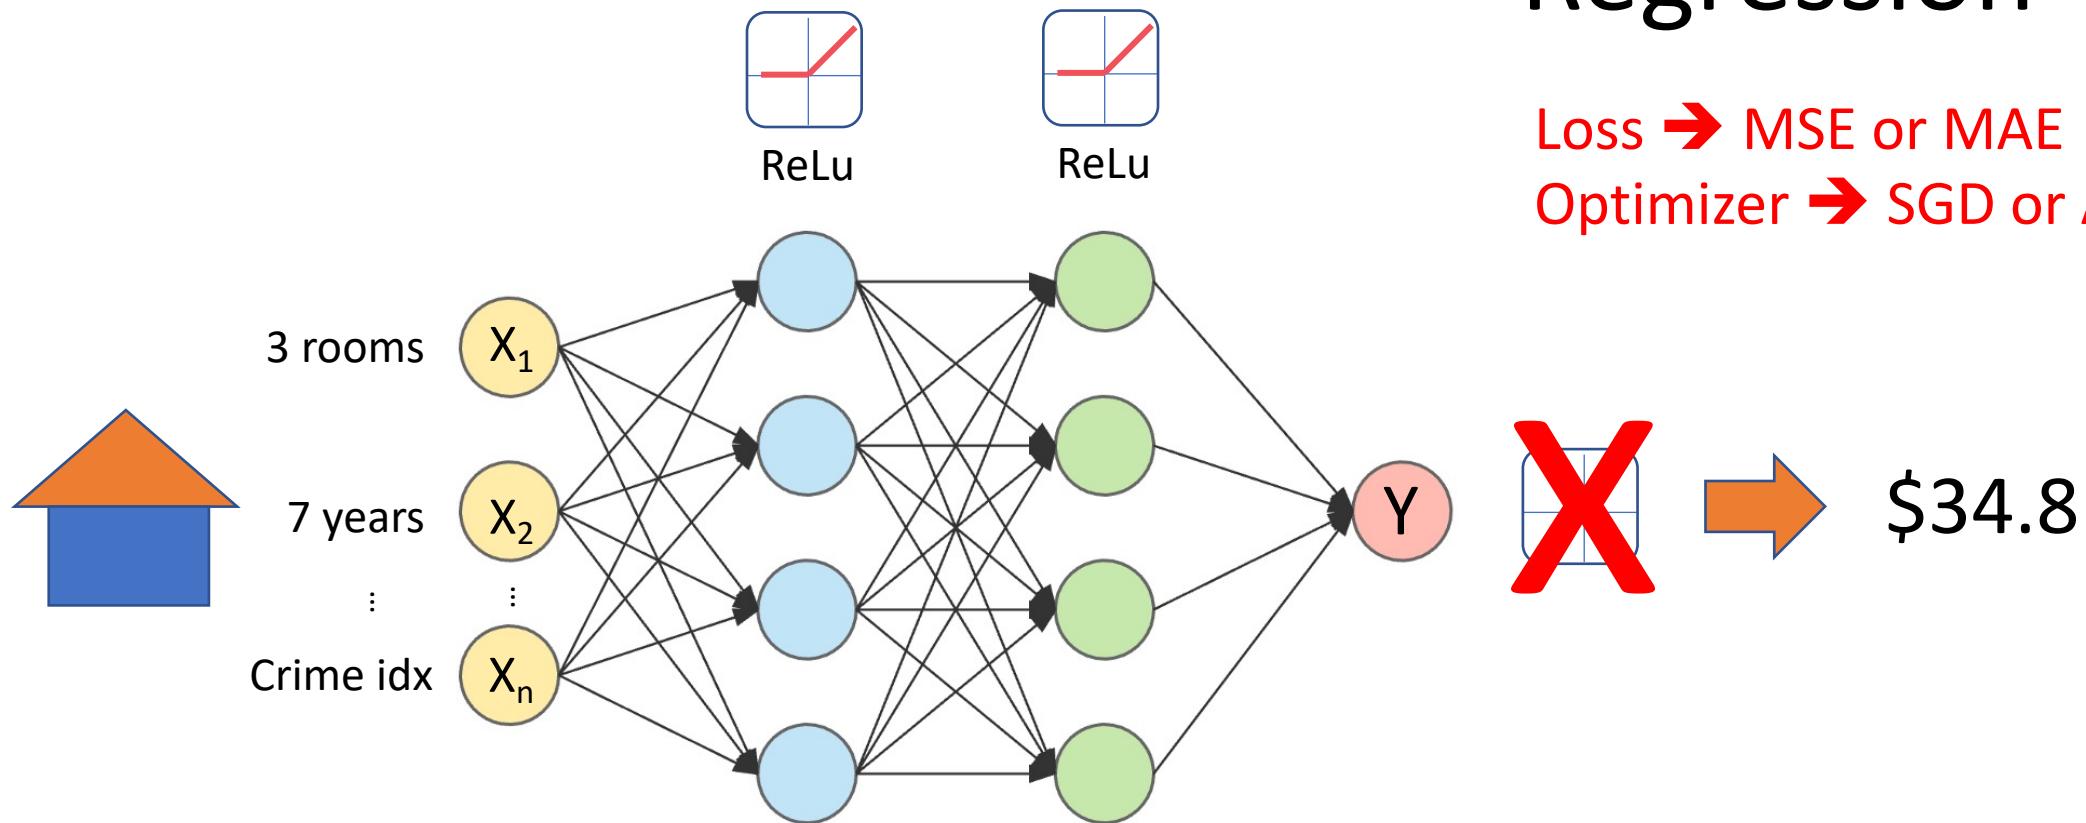


Loss → ?
Optimizer → ?

Regression

Loss → MSE or MAE

Optimizer → SGD or Adam



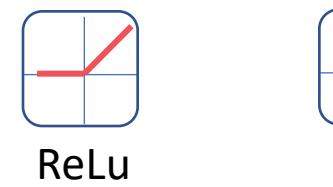
Binary Classification



(28,28)

Flatten

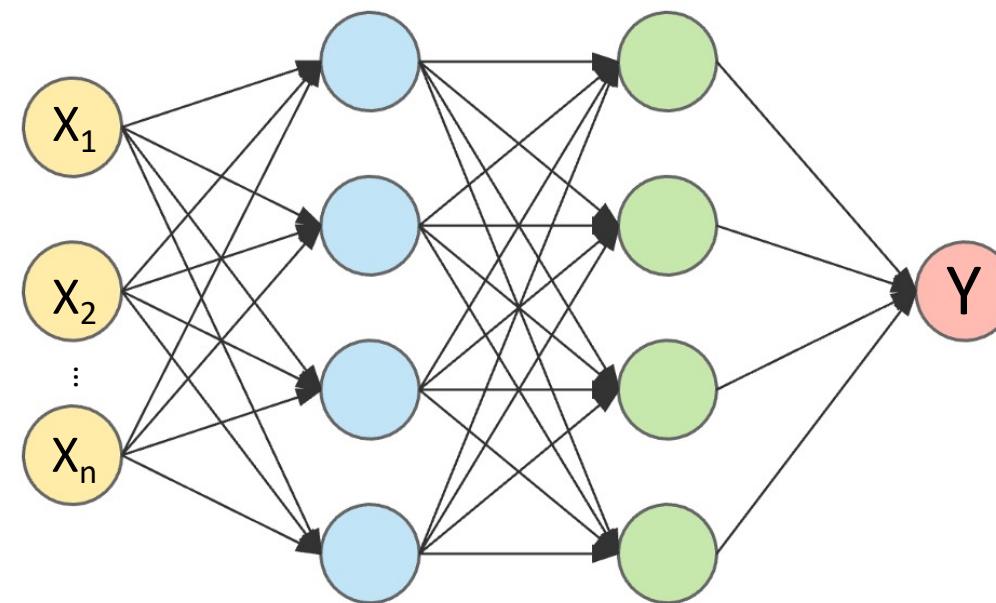
(784)



ReLU



ReLU



Loss → Binary Crossentropy
Optimizer → SGD or Adam



Sigmoid



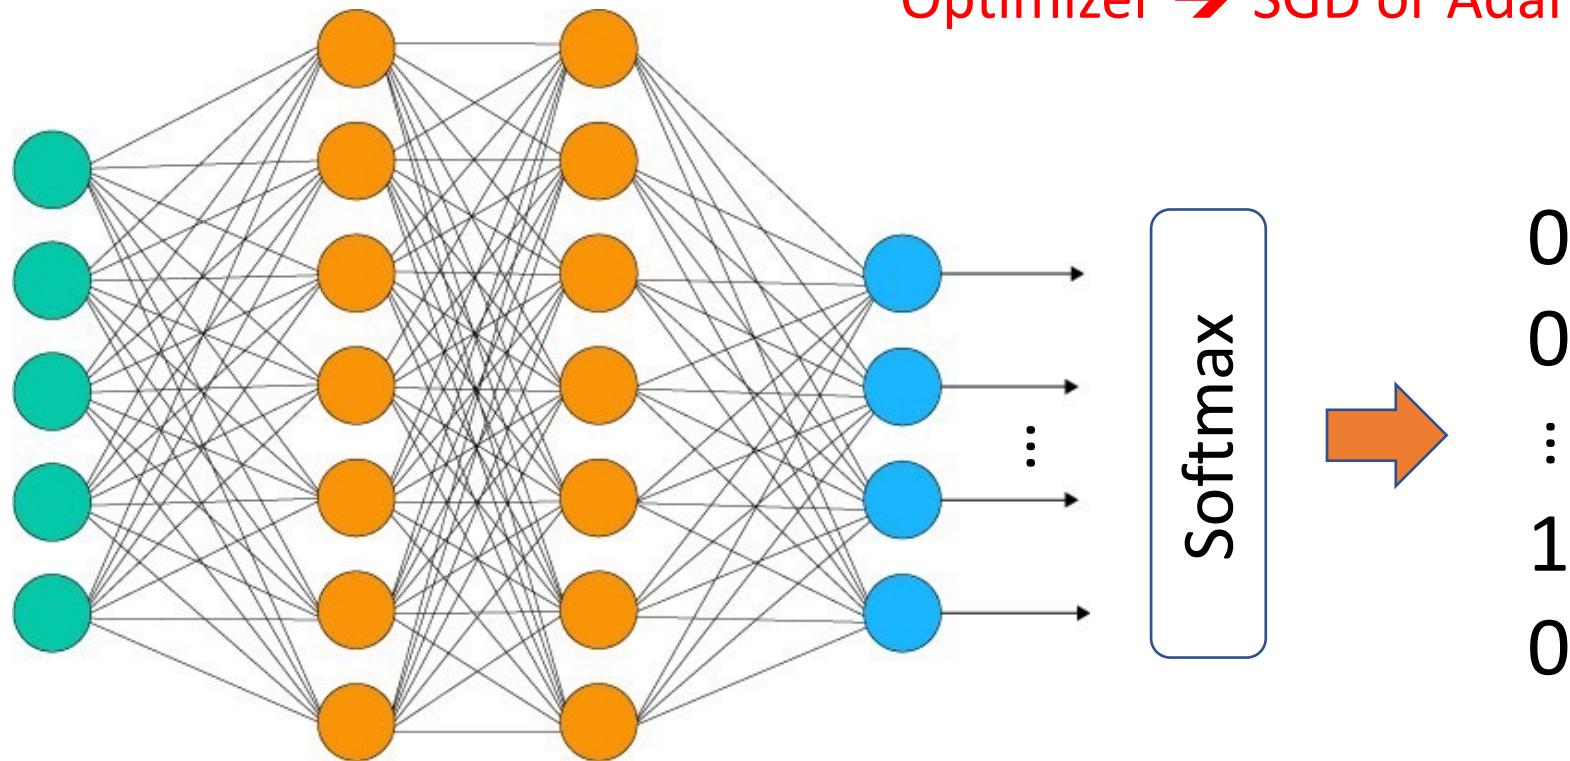
0: Cat
1: Dog

Mult-class Classification

7

Flatten

ReLU
ReLU



Loss → Categorical Crossentropy *
Optimizer → SGD or Adam

* or “Sparse Categorical Crossentropy” if label is 1, 2, 3, ...

Going Further

The Datasets to training and test



Classifying Shoes

Steps to take

1. Get as many examples of shoes as possible
2. Train using these examples
3. Profit!



Steps to take

1. Get as many examples of shoes as possible
2. Train using these examples
3. Profit!

Training accuracy: .920
Training accuracy: .935
Training accuracy: .947
Training accuracy: .961
Training accuracy: .977
Training accuracy: .995
Training accuracy: 1.00

Steps to take

1. ~~Get as many examples of shoes as possible~~
2. ~~Train using these examples~~
3. Profit?



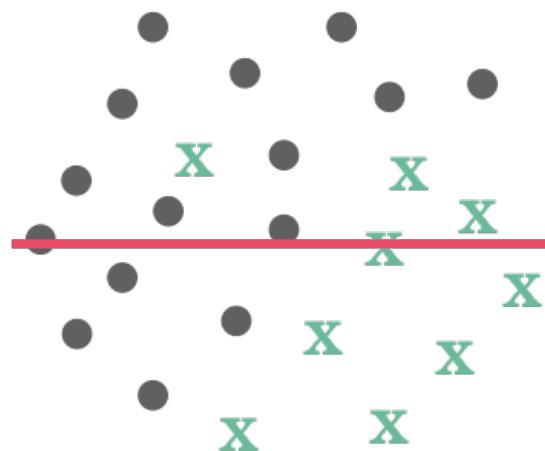
Data

The network ‘sees’ everything. Has no context for measuring how well it does with data it has never previously been exposed to.

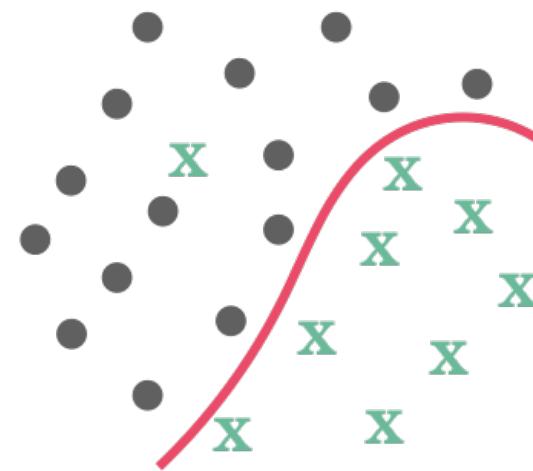
Data

The network ‘sees’ everything. Has no context for measuring how well it does with data it has never previously been exposed to.

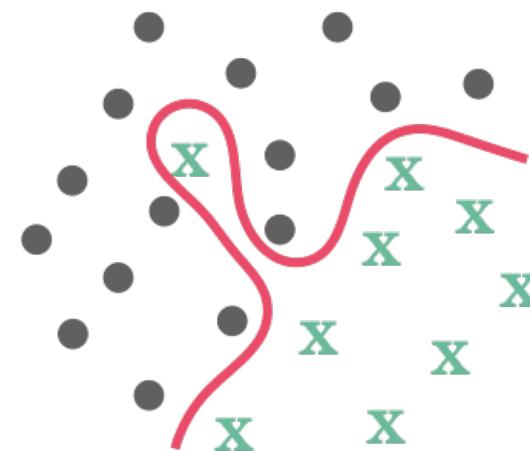
Underfitting



Desired



Overfitting





Data

Validation Data

The network ‘sees’ a subset of your data. You can use the rest to measure its performance against previously unseen data.

Data

Validation Data

Test Data

The network ‘sees’ a subset of your data. You can use an unseen subset to measure its accuracy while training (validation), and then another subset to measure its accuracy after it’s finished training (test).

Is used to evaluate the current training epoch

Is used to evaluate the final model after training

Data

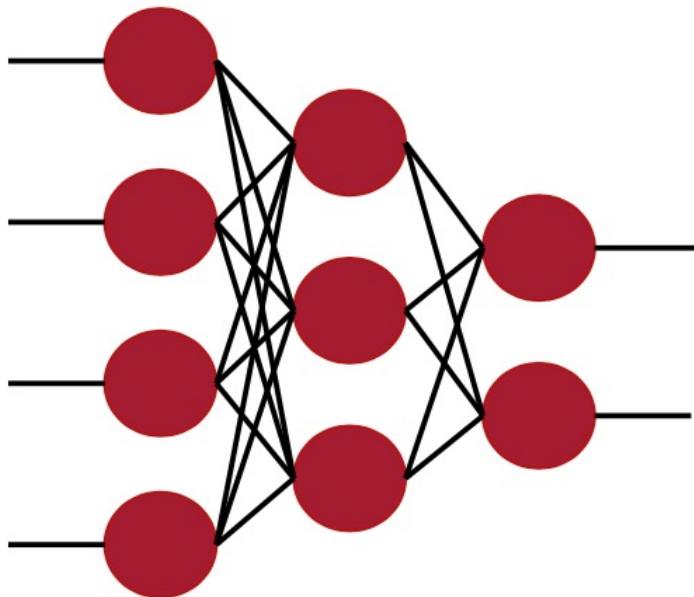
Validation Data

Test Data

Accuracy: 0.999

Accuracy: 0.920

Accuracy: 0.800



Data

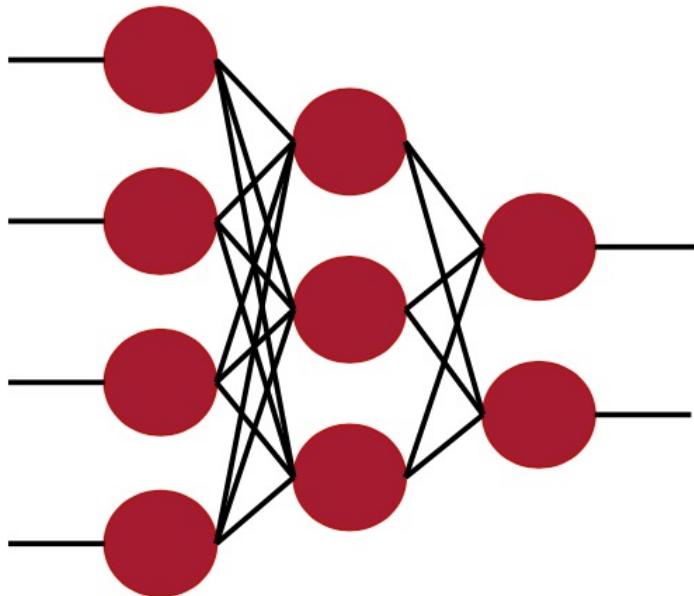
Validation Data

Test Data

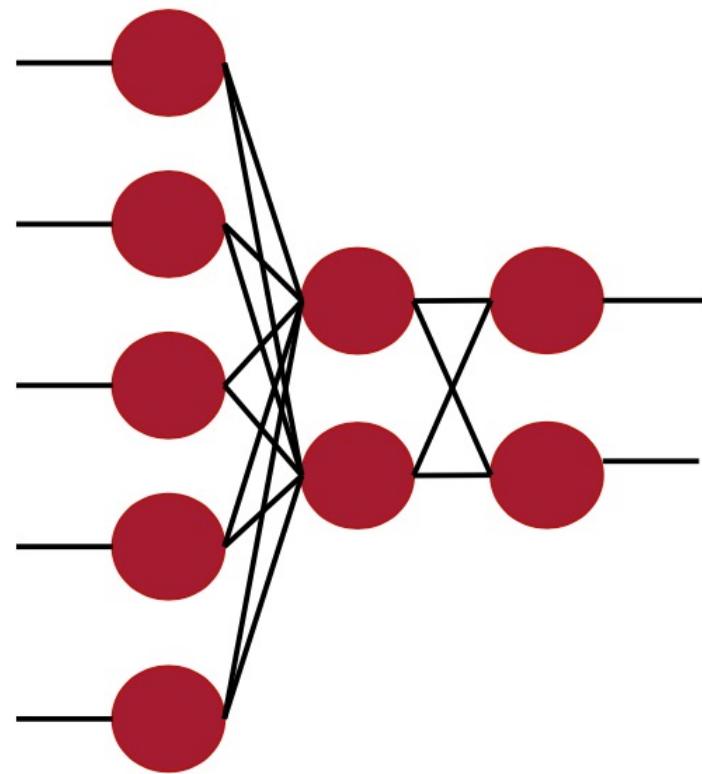
Accuracy: 0.999

Accuracy: 0.920

Accuracy: 0.800



Data



Validation Data

Accuracy: 0.942

Test Data

Accuracy: 0.930

Accuracy: 0.925

```
import tensorflow as tf

data = tf.keras.datasets.mnist
(training_images, training_labels), (val_images, val_labels) = data.load_data()

training_images = training_images / 255.0
val_images = val_images / 255.0

model = tf.keras.models.Sequential(
    [tf.keras.layers.Flatten(input_shape=(28,28)),
     tf.keras.layers.Dense(20, activation=tf.nn.relu),
     tf.keras.layers.Dense(10, activation=tf.nn.softmax)])
```

```
data = tf.keras.datasets.mnist  
  
(tt_images, tt_labels), (val_images, val_labels) = data.load_data()
```

```
test_images = tt_images[:10000]  
test_labels = tt_labels[:10000]
```

```
print(test_images.shape)  
print(test_labels.shape)
```

```
(10000, 28, 28)  
(10000,)
```

```
training_images = tt_images[10000:]  
training_labels = tt_labels[10000:]
```

```
print(training_images.shape)  
print(training_labels.shape)
```

```
(50000, 28, 28)  
(50000,)
```

```
model.fit(training_images, training_labels, epochs=20)
```

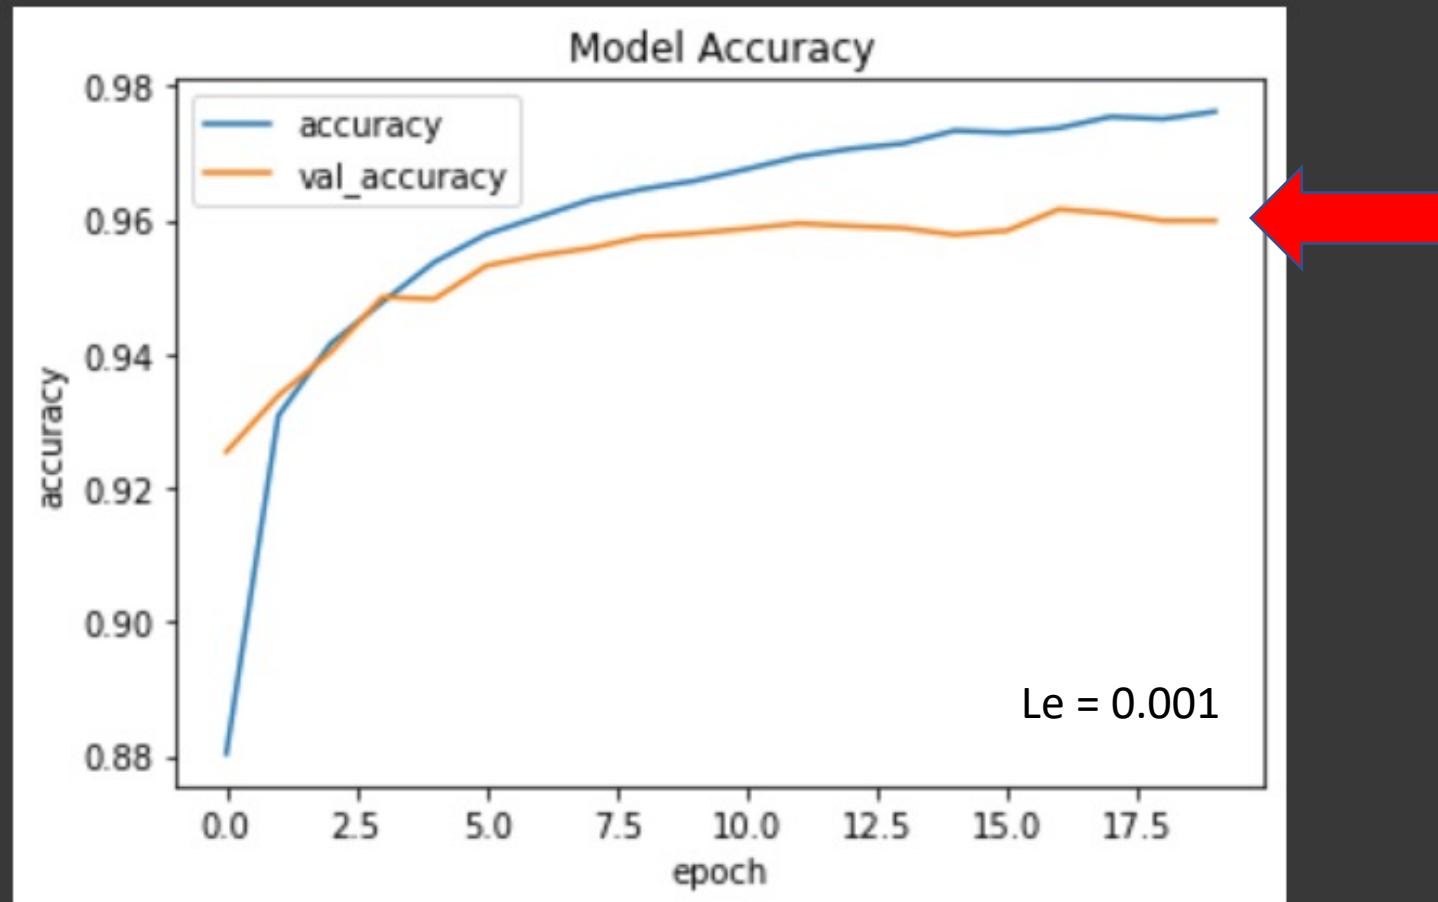
```
Epoch 16/20  
1875/1875 [=====] - 3s 2ms/step - loss: 0.0931 - accuracy: 0.9729  
Epoch 17/20  
1875/1875 [=====] - 3s 2ms/step - loss: 0.0907 - accuracy: 0.9731  
Epoch 18/20  
1875/1875 [=====] - 3s 2ms/step - loss: 0.0892 - accuracy: 0.9735  
Epoch 19/20  
1875/1875 [=====] - 3s 2ms/step - loss: 0.0868 - accuracy: 0.9740  
Epoch 20/20  
1875/1875 [=====] - 3s 2ms/step - loss: 0.0844 - accuracy: 0.9750
```

```
model.fit(training_images, training_labels,  
          validation_data=(val_images, val_labels),  
          epochs=20)
```

```
model.fit(training_images, training_labels,  
          validation_data=(val_images, val_labels),  
          epochs=20)
```

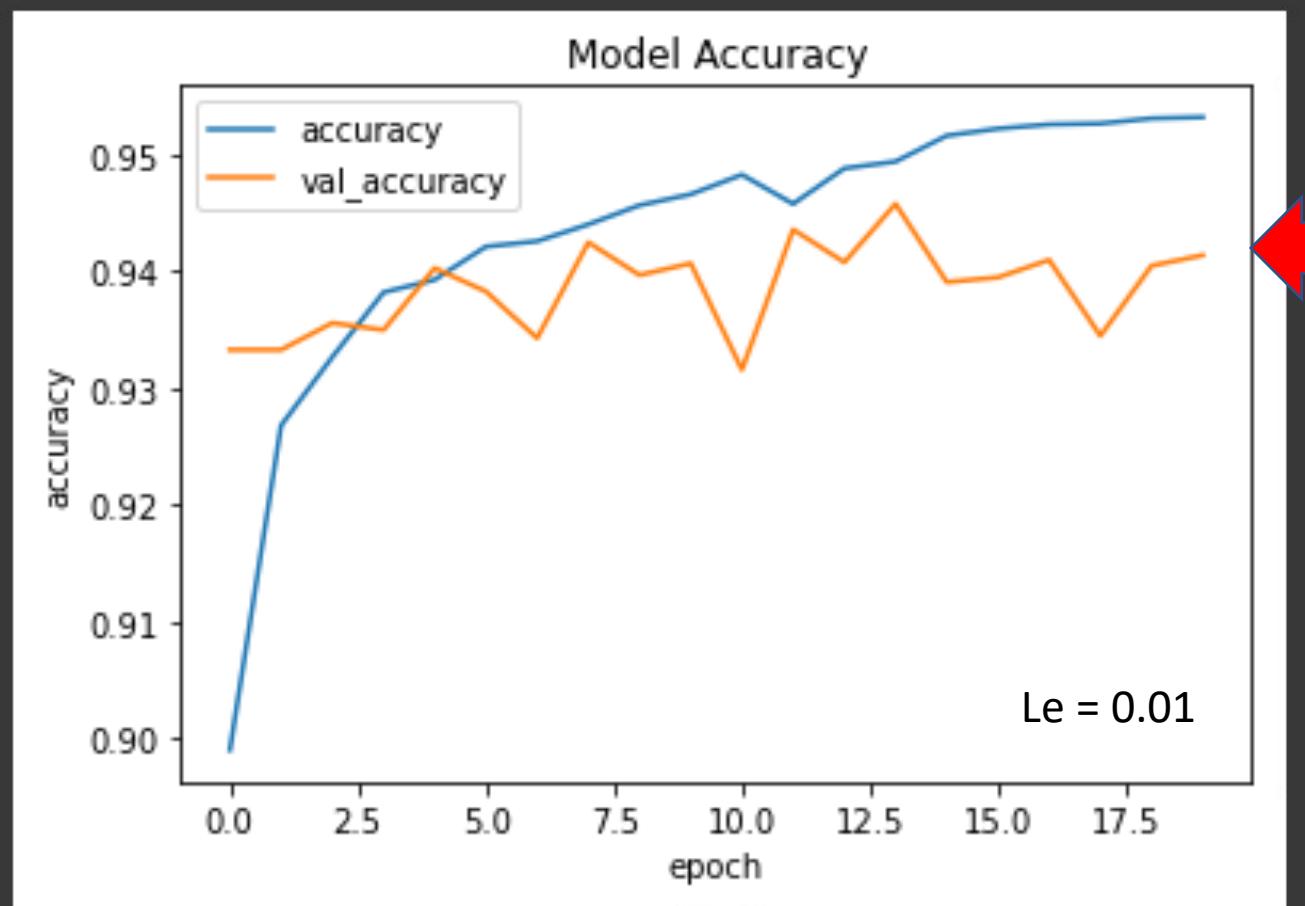
```
Epoch 14/20  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0960 - accuracy: 0.9713 - val_loss: 0.1410 - val_accuracy: 0.9588  
Epoch 15/20  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0933 - accuracy: 0.9722 - val_loss: 0.1353 - val_accuracy: 0.9626  
Epoch 16/20  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0910 - accuracy: 0.9729 - val_loss: 0.1356 - val_accuracy: 0.9622  
Epoch 17/20  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0880 - accuracy: 0.9734 - val_loss: 0.1339 - val_accuracy: 0.9625  
Epoch 18/20  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0855 - accuracy: 0.9740 - val_loss: 0.1349 - val_accuracy: 0.9619  
Epoch 19/20  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0832 - accuracy: 0.9754 - val_loss: 0.1341 - val_accuracy: 0.9624  
Epoch 20/20  
1875/1875 [=====] - 4s 2ms/step - loss: 0.0808 - accuracy: 0.9759 - val_loss: 0.1340 - val_accuracy: 0.9626
```

```
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label='val_accuracy')
plt.title('Model Accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(loc='upper left')
plt.show()
```



If validation accuracy goes down, even if train accuracy goes up, means that probably the model is overfitting. In this case the training (epochs) should terminate.

```
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label='val_accuracy')
plt.title('Model Accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(loc='upper left')
plt.show()
```



If validation accuracy seems “instable”, could be that Learning Rate is high (try to reduce it).

```
model.evaluate(test_images, test_labels)
```

```
313/313 [=====] - 1s 2ms/step - loss: 0.1495 - accuracy: 0.9569
```

Data

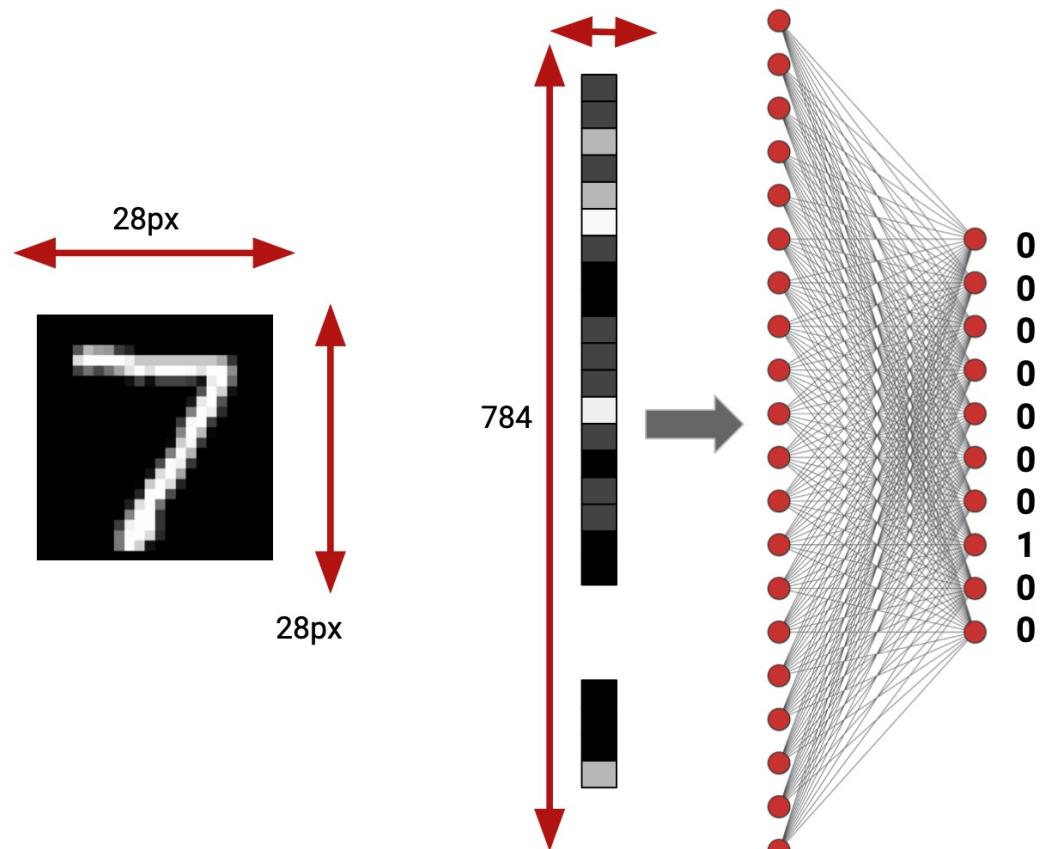
Validation Data

Test Data

Accuracy: 0.976

Accuracy: 0.963

Accuracy: 0.957



In summary

Training Data → Used to train **model parameters**

Validation Data → Used to determine what **model hyperparameters** to adjust (and re-training)

Test Data → Used to get **model final performance metric**

Digits Classification: validation and test dataset

Code Time!

TF_MNIST_Classification_v2.ipynb



Going Further

Classification Model Performance Metrics



Class = [1]

actual = [1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0]



Class = [0]

prediction = [0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 0, 1]



Model Performance (Confusion Matrix)

		predicted condition	
		Cat	Dog
		6	2
true condition	Cat	6	2
	Dog	1	3

Model Performance (Confusion Matrix)

		predicted condition	
		Cat	Dog
		True Positive (TP)	False Negative (FN) (type II error)
true condition	Cat	6	2
	Dog	1	3
		False Positive (FP) (Type I error)	True Negative (TN)

Model Performance (Confusion Matrix)

		predicted condition	
total population (P + N)		prediction positive (PP)	prediction negative (PN)
true condition	condition positive (P)	True Positive (TP)	False Negative (FN) (type II error)
	condition negative (N)	False Positive (FP) (Type I error)	True Negative (TN)

Type I error (false positive)



Type II error (false negative)

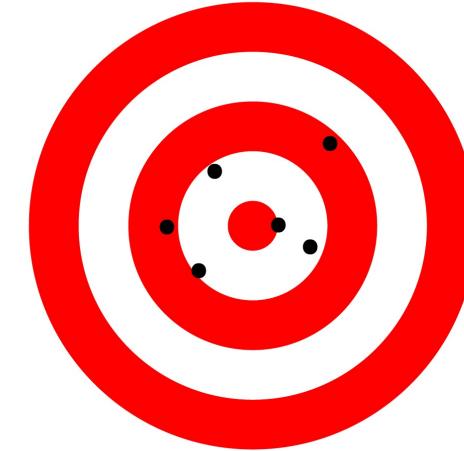


Precision vs. Accuracy

In a set of measurements, **accuracy** is closeness of the measurements to a specific value, while **precision** is the closeness of the measurements to each other.



High Precision, High Accuracy



Low Precision, High Accuracy



High Precision, Low Accuracy



Low Precision, Low Accuracy

Accuracy , Precision and Recall

$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{(\text{P} + \text{N})} = \frac{\text{TP} + \text{TN}}{(\text{TP} + \text{TN} + \text{FP} + \text{FN})} = \frac{6 + 3}{(6 + 3 + 1 + 2)} = \frac{9}{12} = 0.75$$

$$\text{Precision} = \frac{\text{TP}}{(\text{TP} + \text{FP})} = \frac{6}{(6 + 1)} = \frac{6}{7} = 0.86$$

Total Positive
Total Predict Positive

$$\text{Recall} = \frac{\text{TP}}{(\text{TP} + \text{FN})} = \frac{6}{(6 + 2)} = \frac{6}{8} = 0.75$$

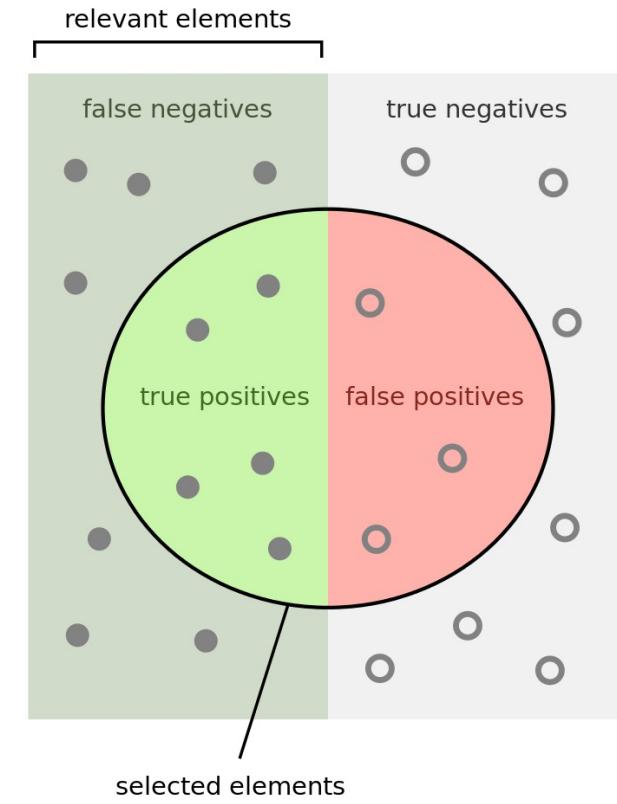
Total Positive
Total Actual Positive

F1-Score

$$F1 = \frac{2 \times (\text{Precision} * \text{Recall})}{(\text{Precision} + \text{Recall})}$$

$$F1 = \frac{2 \times (0.86 * 0.75)}{(0.86 + 0.75)} = \frac{2 \times 0.65}{1.61} = 0.80$$

The F1-score is a way of combining the precision and recall of the model



How many selected items are relevant?

$$\text{Precision} = \frac{\text{green}}{\text{green} + \text{red}}$$

How many relevant items are selected?

$$\text{Recall} = \frac{\text{green}}{\text{green} + \text{grey}}$$

```
1 from sklearn.metrics import classification_report
```

```
1 actual = [1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0]  
2 prediction = [0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 0, 1]
```

```
1 target_names = ['Dogs', 'Cats']
```

```
1 print(classification_report(actual, prediction, target_names=target_names))
```

	precision	recall	f1-score	support
Dogs	0.60	0.75	0.67	4
Cats	0.86	0.75	0.80	8
accuracy			0.75	12
macro avg	0.73	0.75	0.73	12
weighted avg	0.77	0.75	0.76	12

Classification Report

Code Time!

Classification_Report.ipynb



Reading Material

Main references

- [Harvard School of Engineering and Applied Sciences - CS249r: Tiny Machine Learning](#)
- [Professional Certificate in Tiny Machine Learning \(TinyML\) – edX/Harvard](#)
- [Introduction to Embedded Machine Learning \(Coursera\)](#)
- [Text Book: "TinyML" by Pete Warden, Daniel Situnayake](#)

I want to thank [Laurence Moroney](#) from Google, Harvard professor [Vijay Janapa Reddi](#), Ph.D. student [Brian Plancher](#) and their staff for preparing the excellent material on TinyML that is the basis of this course at UNIFEI.

The IESTI01 course is part of the [TinyML4D](#), an initiative to make TinyML education available to everyone globally.

Thanks
And stay safe!

