

IESTI01 - TinyML

Part 2 - TinyML Applications

Prof. Marcelo Rovai

June 9th, 2021



Tiny Machine Learning (TinyML)

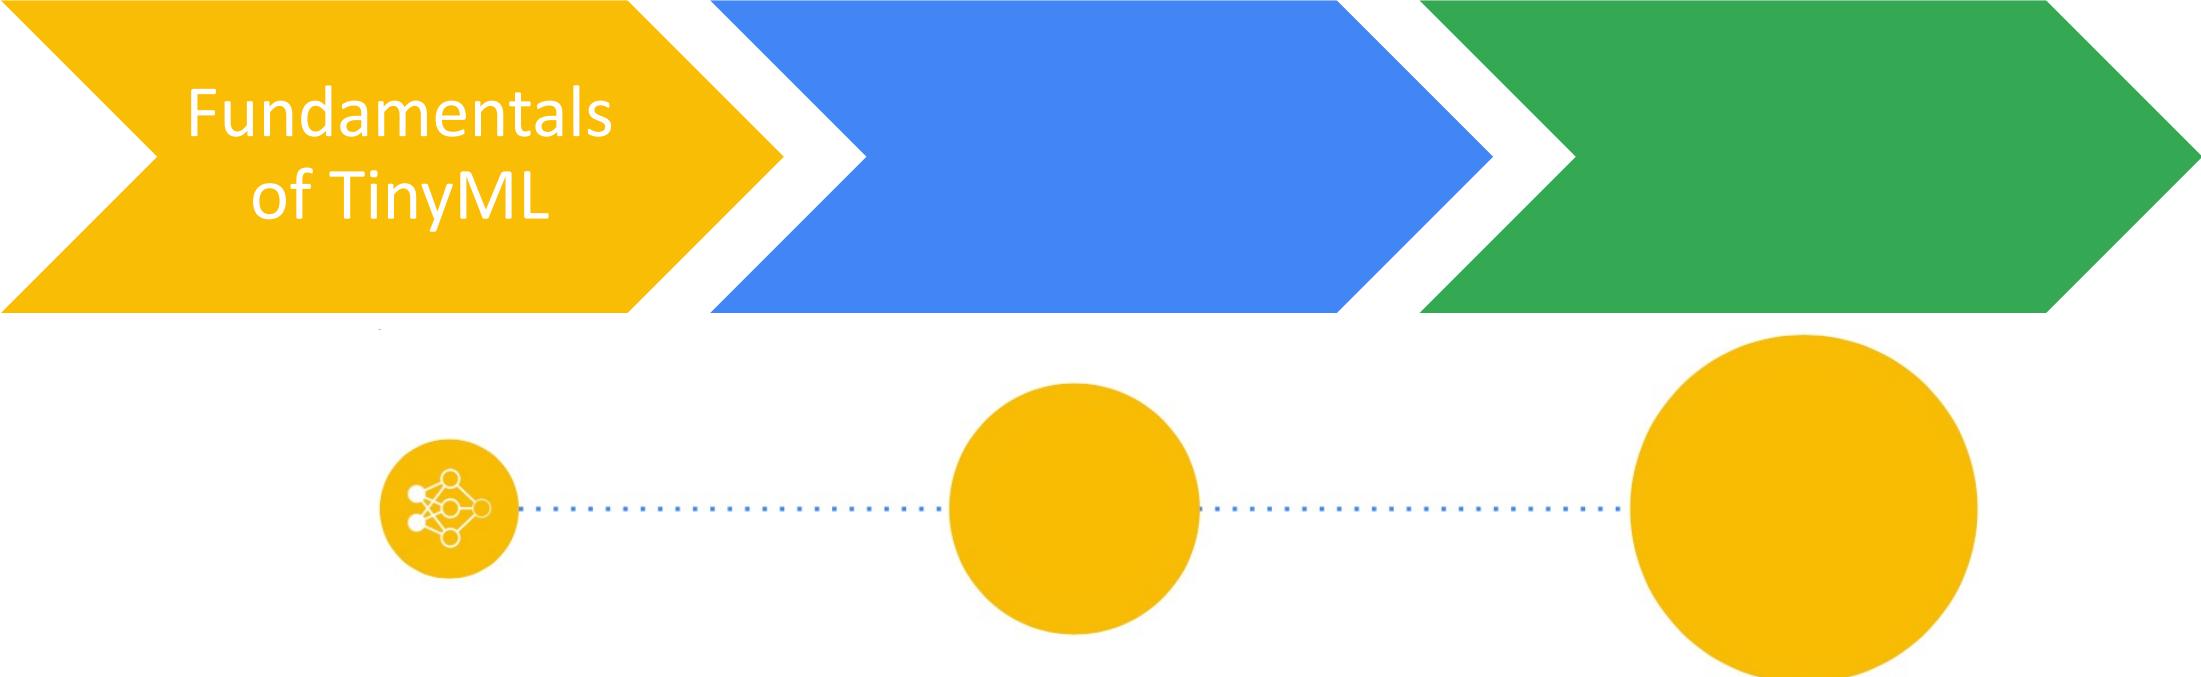
What we learned so far

What is Tiny Machine Learning (**TinyML**)?

- Fast-growing field of **machine learning**
- Algorithms, **hardware, and software**
- **On-device** sensor data analytics
- Extreme **low power** consumption
- **Always-on ML** use-cases
- **Battery**-operated devices

What we already learned?

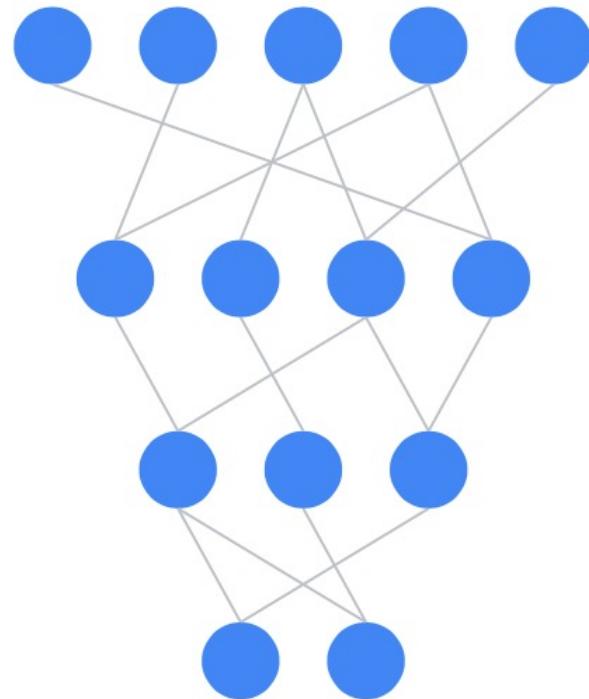
April/May



Fundamentals
of TinyML

In April and May (Part 1) we introduced ML with TensorFlow. Was all about talking about what is the **language of machine learning**.

Total Recall from Part 1



“Language” for Part 1

Neural Network

Training

Validation Data

Gradient Descent

Inference

Test Data

Loss Function

Features

Classification

Filters

Overfitting

Regression

Kernels

Data augmentation

Responsible AI

CNNs

DNNs

Preprocessing

“Language” for Part 1

Neural Network

Training

Training Data

Gradient Descent

Inference

Validation Data

Loss Function

Features

Classification

Kernels

Filters

Overfitting

Regression

CNNs

Data augmentation

Responsible AI

DNNs

Preprocessing

“Language” for Part 1

Neural Network

Training

Training Data

Gradient Descent

Inference

Validation Data

Loss Function

Features

Classification

Kernels

Filters

Overfitting

Regression

CNNs

DNNs

Data augmentation

Responsible AI

Preprocessing

“Language” for Part 1

Neural Network

Training

Training Data

Gradient Descent

Inference

Validation Data

Loss Function

Features

Classification

Kernels

Filters

Overfitting

Regression

CNNs

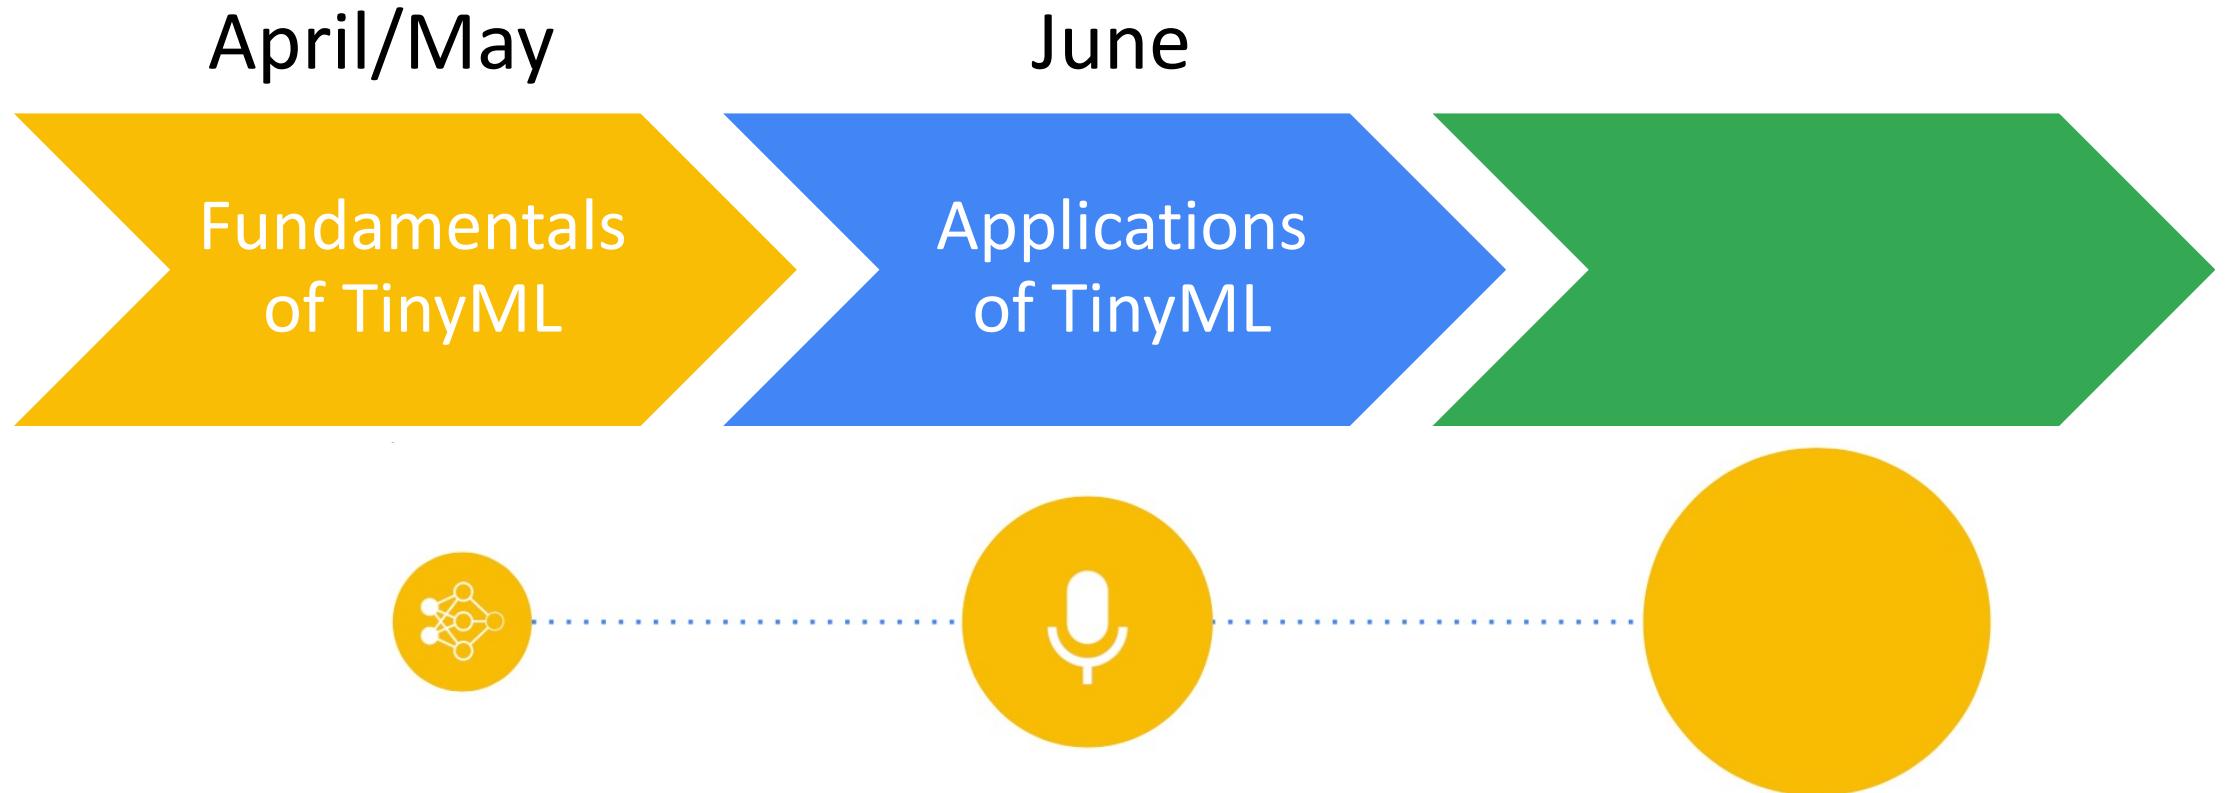
Data augmentation

Responsible AI

DNNs

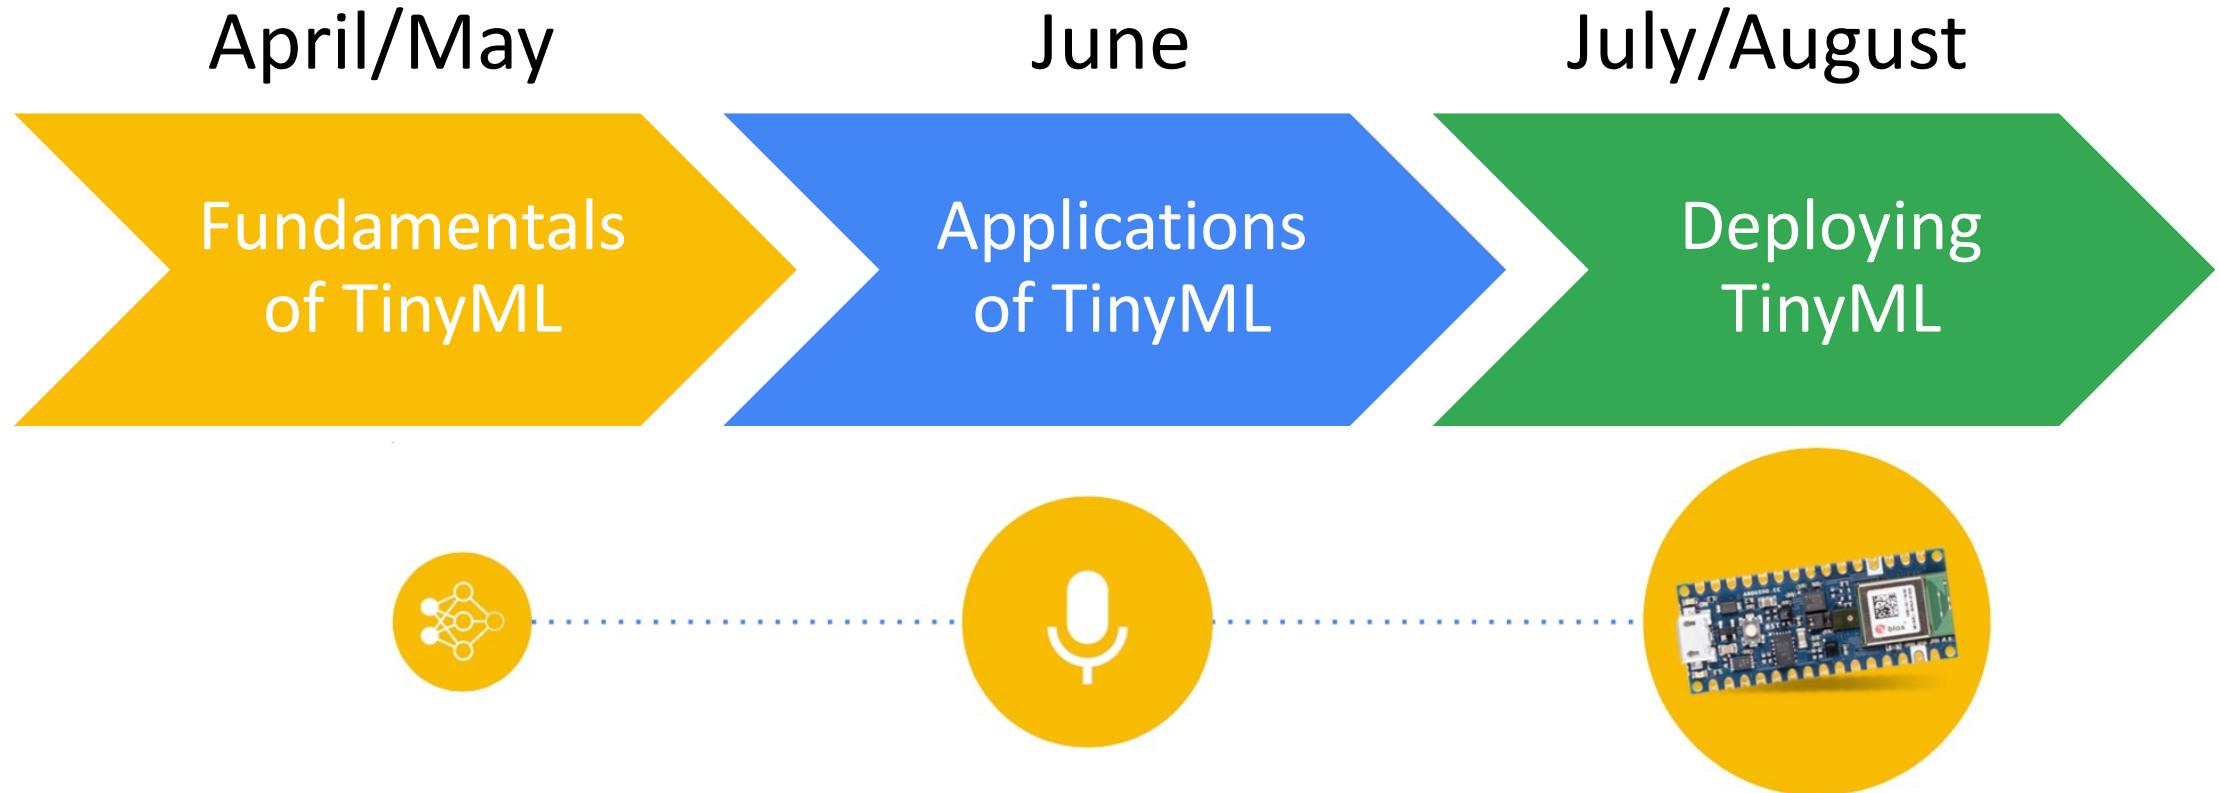
Preprocessing

What we will learn?



In June (Part 2), We will learn the importance of dataset engineering, get a sneak peek into different TinyML applications, as keyword spotting (“Alexa”), gesture recognition, and image detection.

What we will learn?



In July and August (Part 3), We will learn how to deploy models on a real microcontroller. Along the way we will explore the challenges unique to and amplified by TinyML (e.g., preprocessing, post-processing, dealing with resource constraints).



TensorFlow



TensorFlow Lite

Train a model

Convert
model

Optimize
model

Deploy
model at
Edge

Make
inferences
at Edge





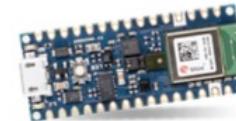
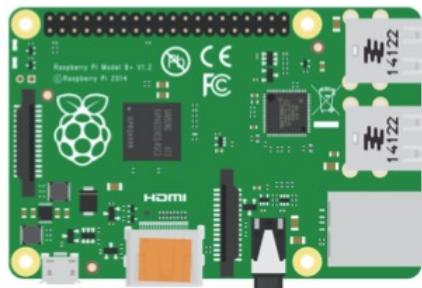
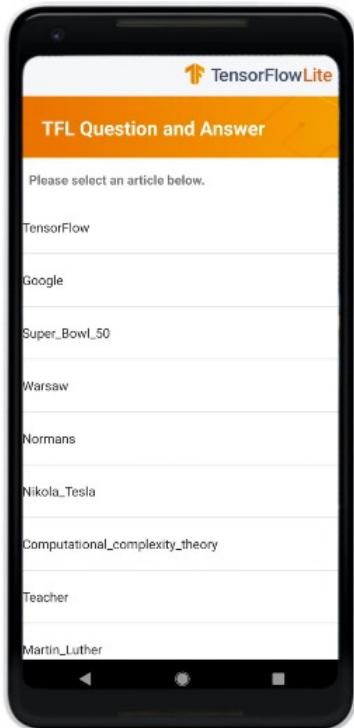
Train a model

Convert
model

Optimize
model

Deploy
model at
Edge

Make
inferences
at Edge

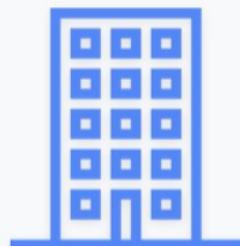


Tiny Machine Learning (TinyML) Applications

TinyML Application Areas



Home



Office

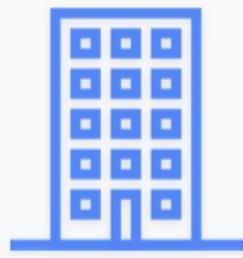


Industry

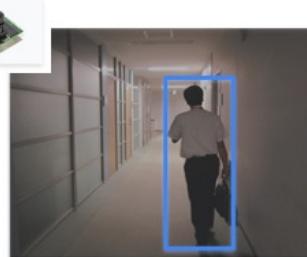
TinyML Application Areas



Home



Office



Industry



Questions

- How do we **capture** the data to feed into the neural network?
- How do you **design** the neural network to take in the speech signal?
- What **dataset** does the neural network need to be trained?
- How do we **pre-process** the data for neural network inference?
- How do you **post-process** the neural network output?
- How do you make sure there is no **bias** in the dataset?
- How do you **deploy** this on the microcontroller?

Endpoints Have Sensors, Tons of Sensors

Motion Sensors

Gyroscope, Radar,
Accelerometer

Acoustic Sensors

Ultrasonic, Microphones,
Geophones, Vibrometers

Environmental Sensors

Temperature, Humidity,
Pressure, IR, etc.

Touchscreen Sensors

Capacitive, IR

Image Sensors

Thermal, Image

Biometric Sensors

Fingerprint, Heart rate, etc.

Force Sensors

Pressure, Strain

Rotation Sensors

Encoders

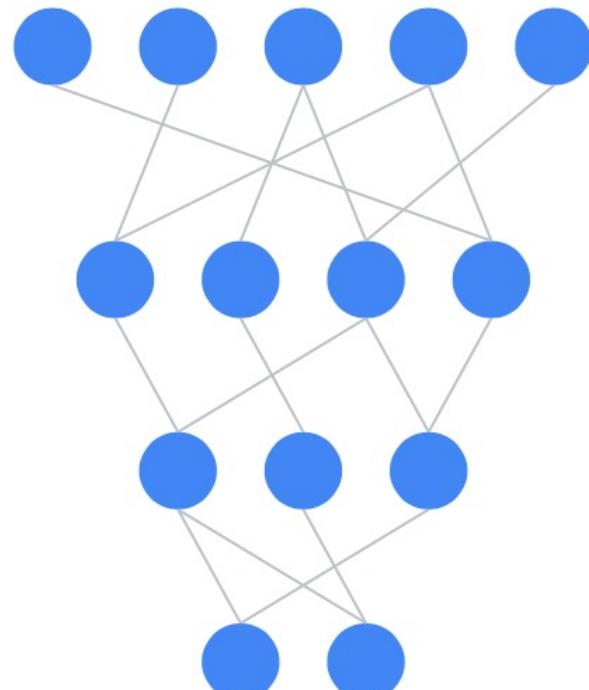
Sensors

Acoustic Sensors
Ultrasonic, Microphones,
Geophones, Vibrometers

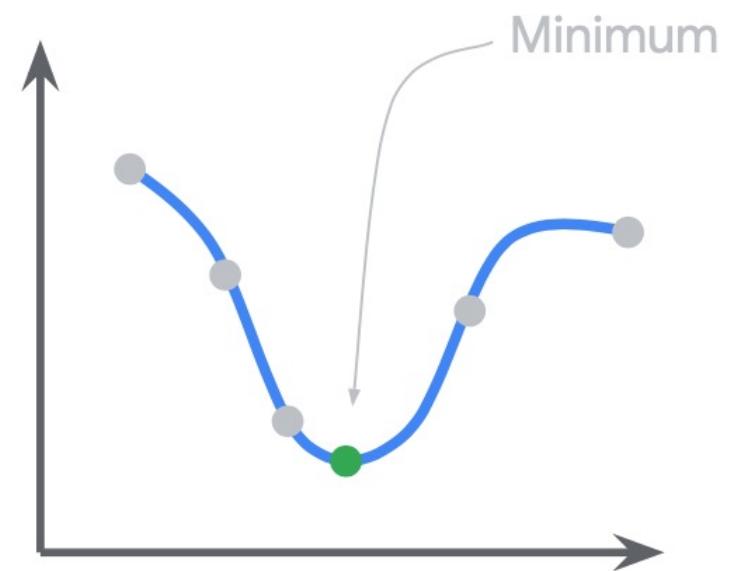
Image Sensors
Thermal, Image

Motion Sensors
Gyroscope, Radar,
Accelerometer

Models



Metrics



End-to-end **TinyML** application design

Datasets Preprocessing

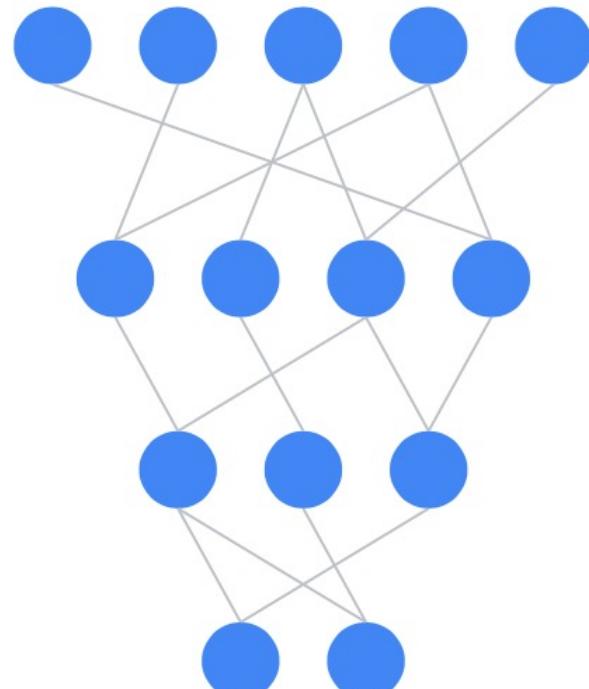
Quantization Pruning

Resource constraints

Sound

Vision

Vibration

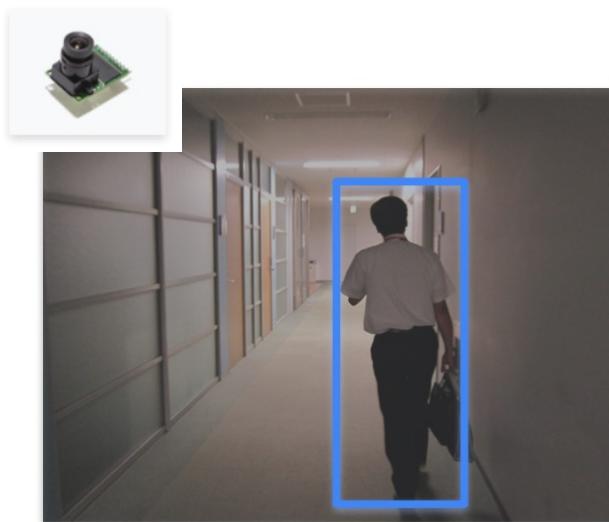


End-to-end **TinyML** application design

Sound



Vision



Vibration



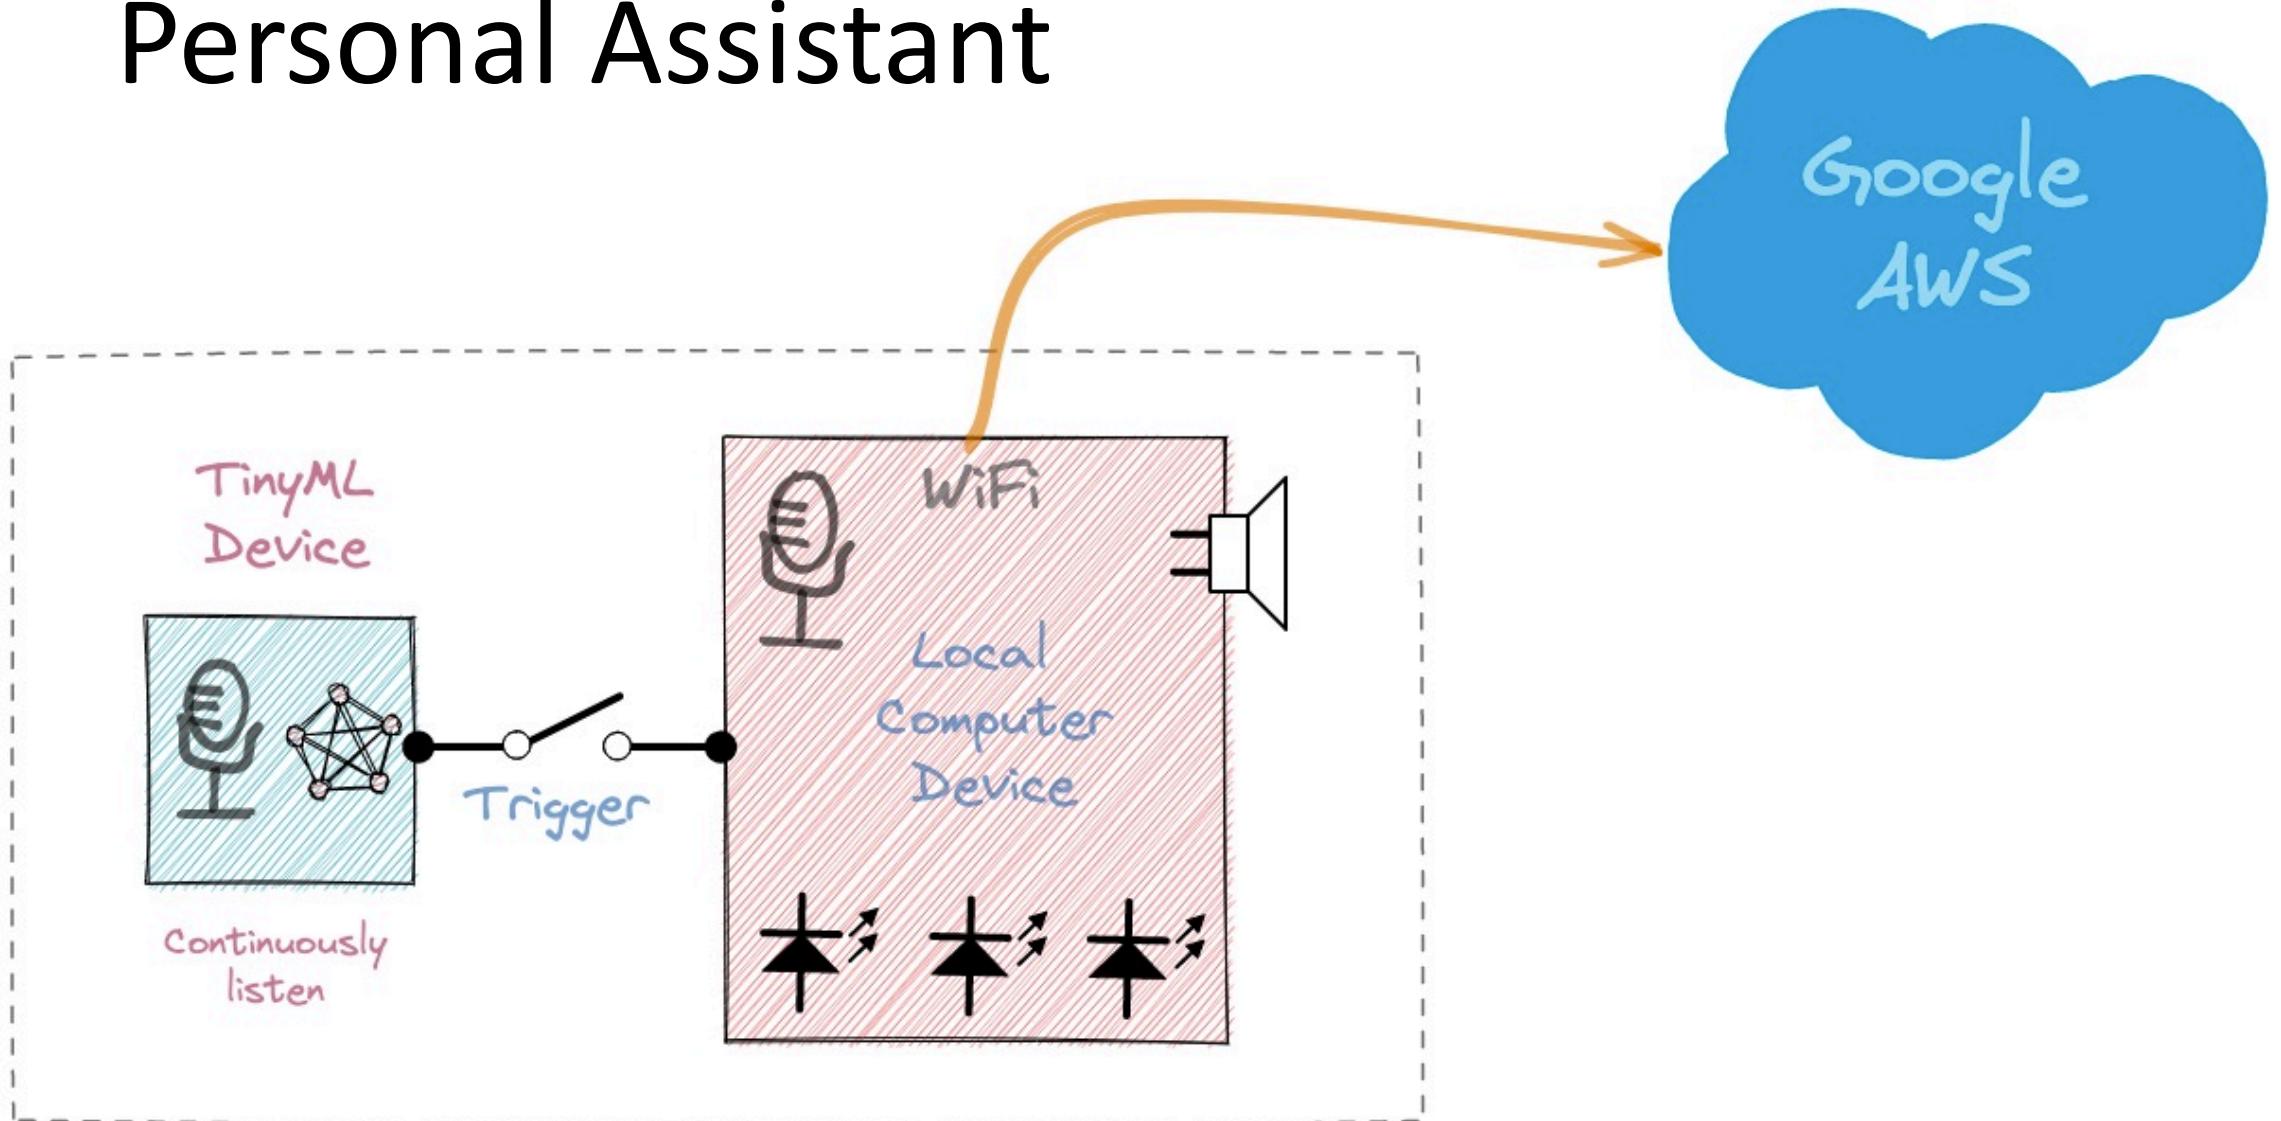
TinyML Application

Example

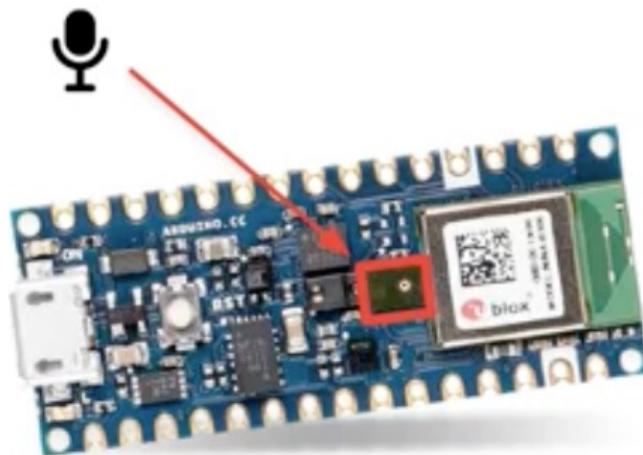
Personal Assistant



Personal Assistant



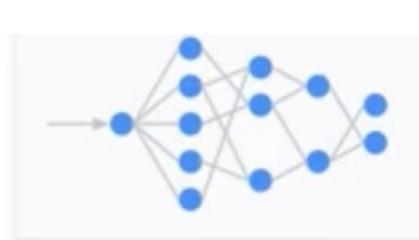
“Cascade” Detection: multi-stage model



- 1 Continuously listen on the microcontroller

2

- Process the data with **TinyML** at the edge



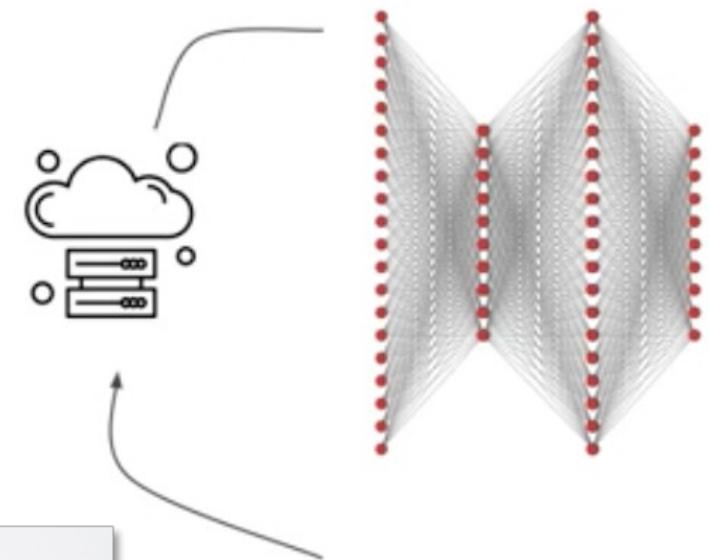
3

- Process on a secondary larger model on a larger local device



5

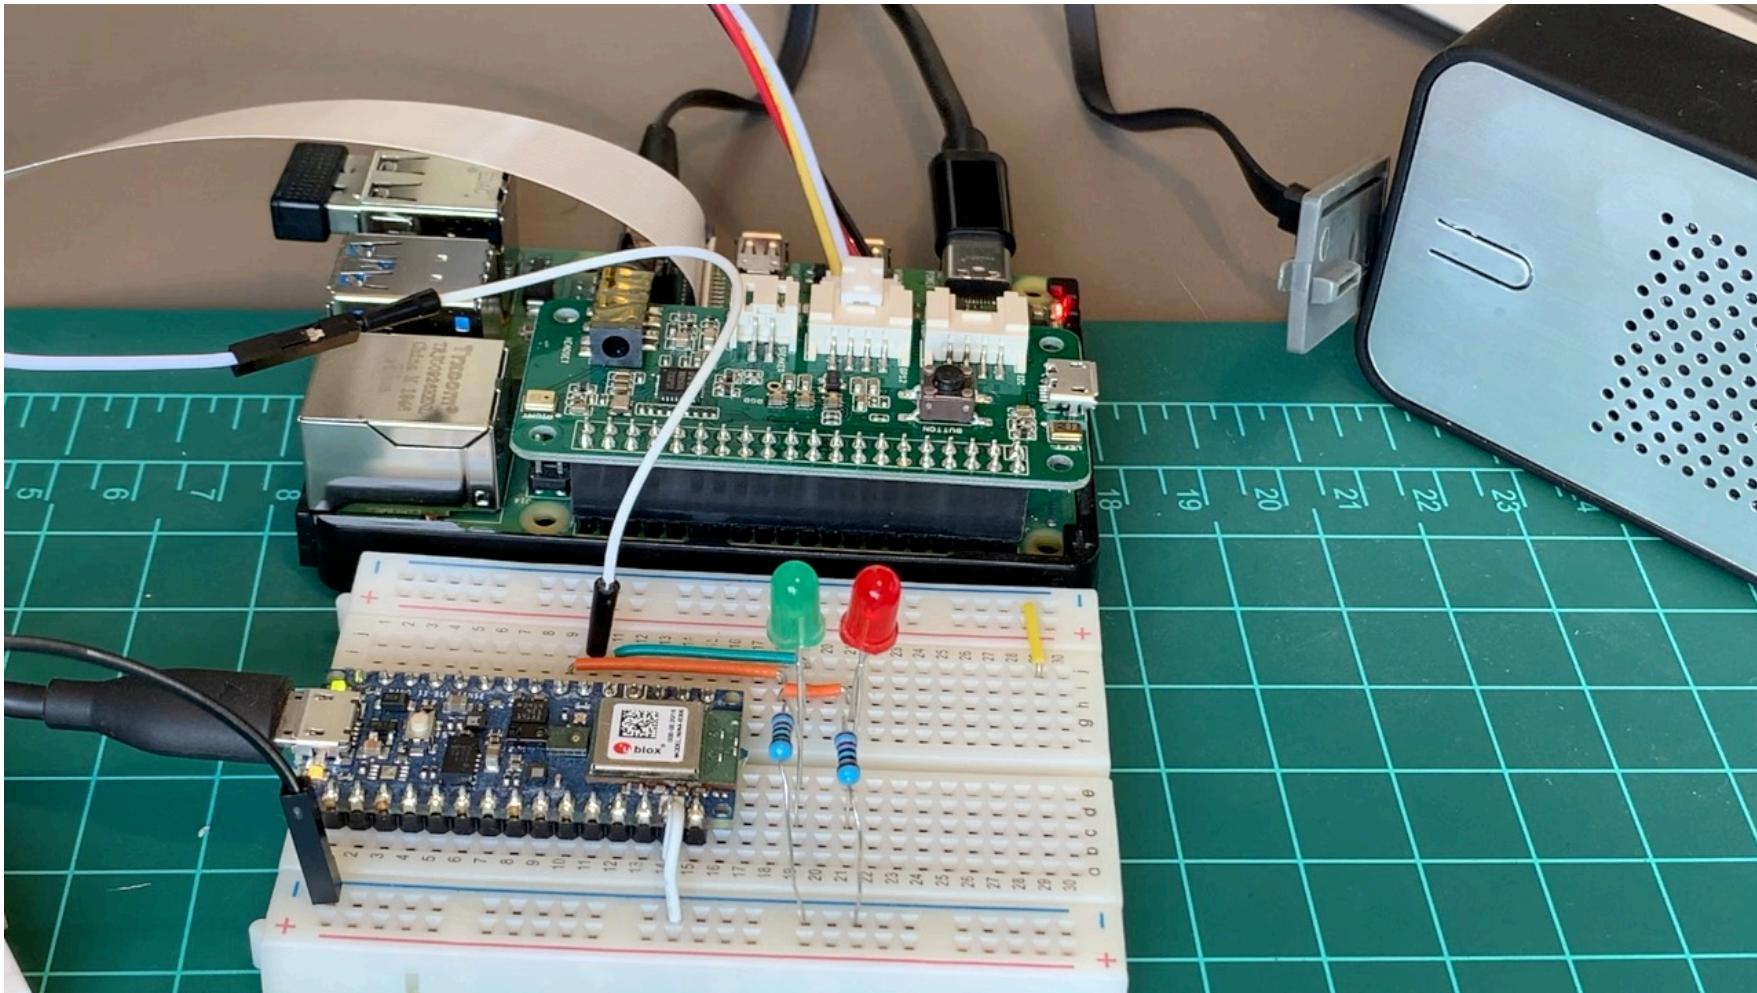
- Process the full speech data with a large model in the cloud



4

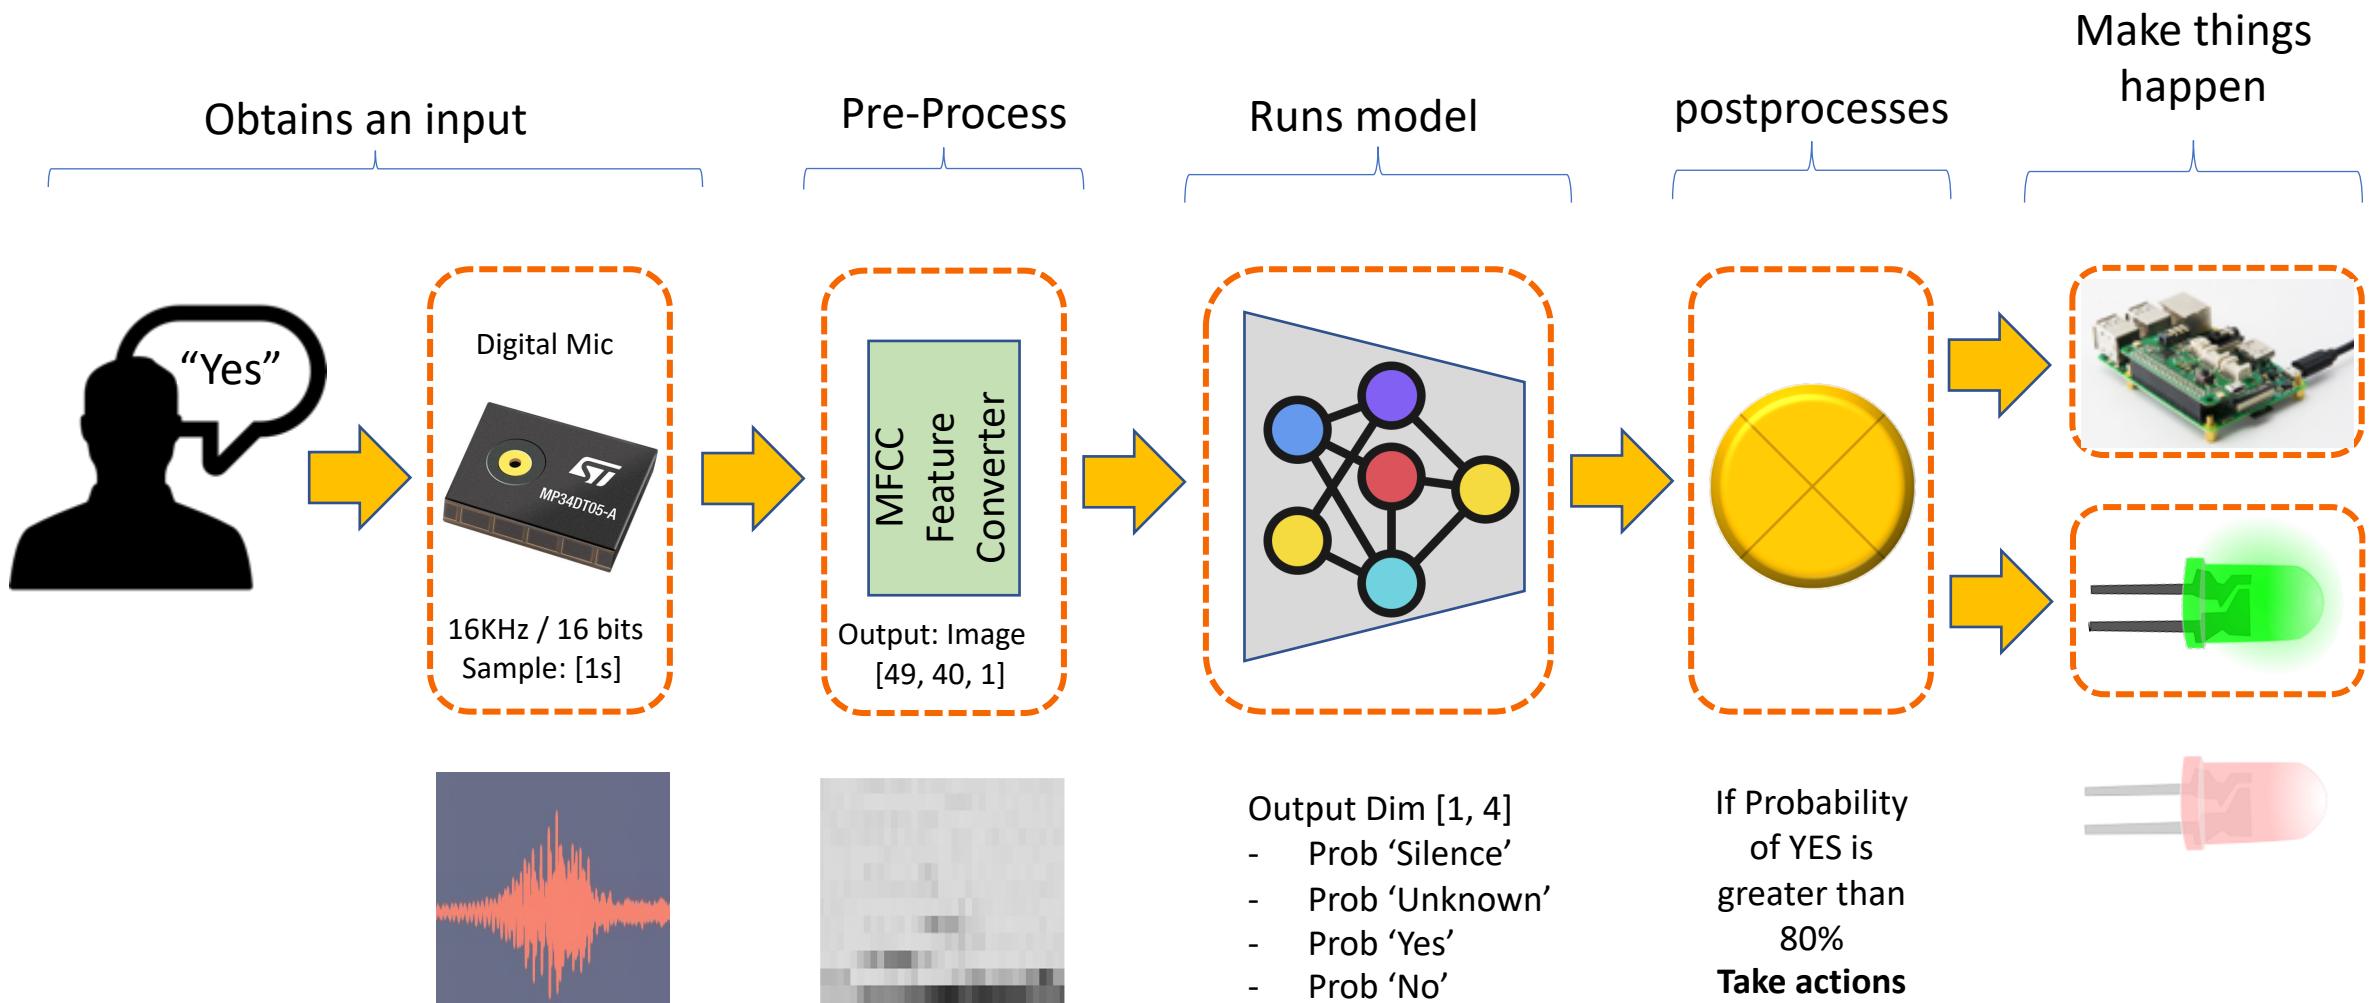
- Send the data to the cloud when triggered

KeyWord Spotting (KWS)

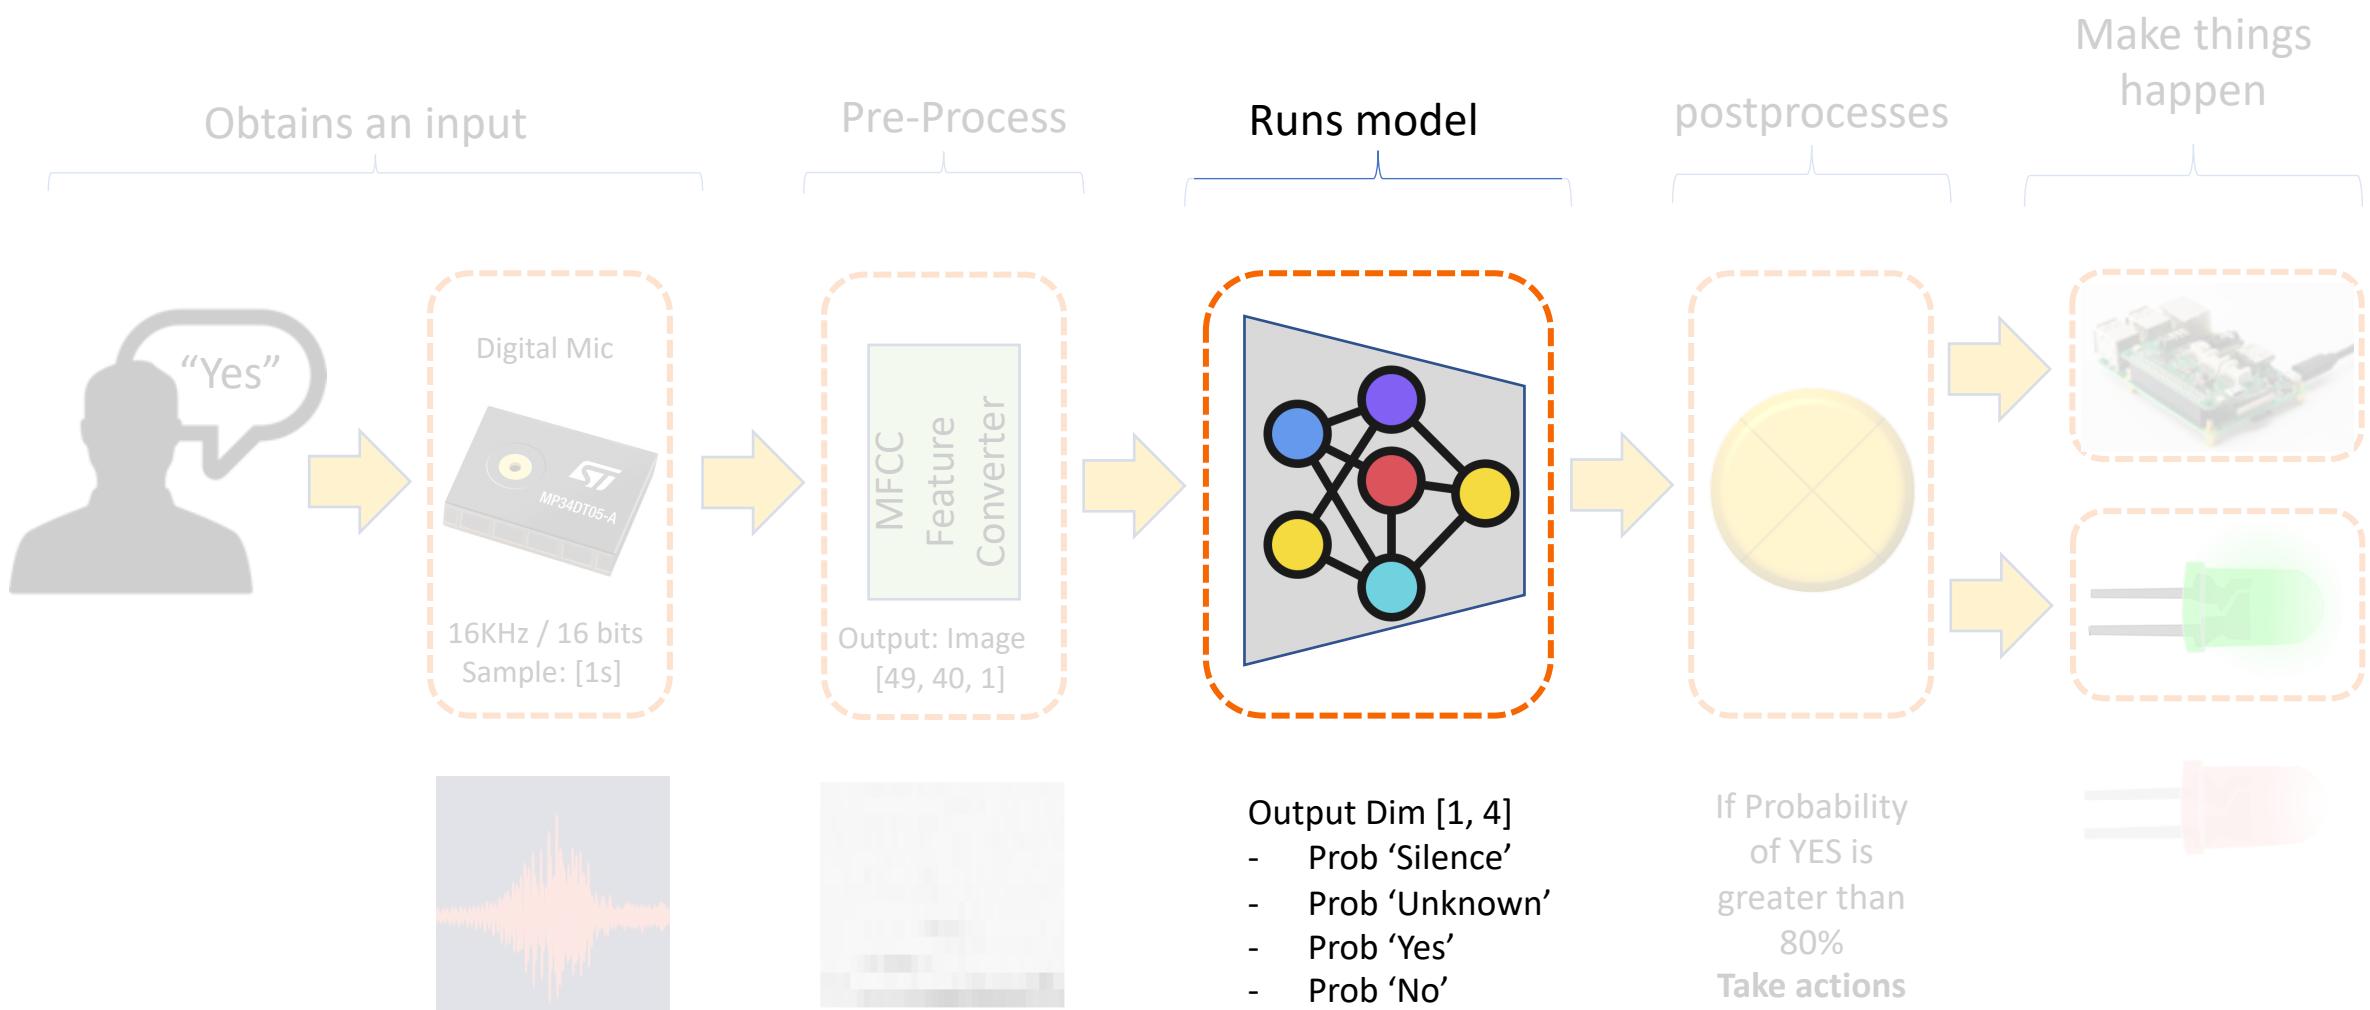


<https://mjrobot.org/2021/01/27/building-an-intelligent-voice-assistant-from-scratch/>

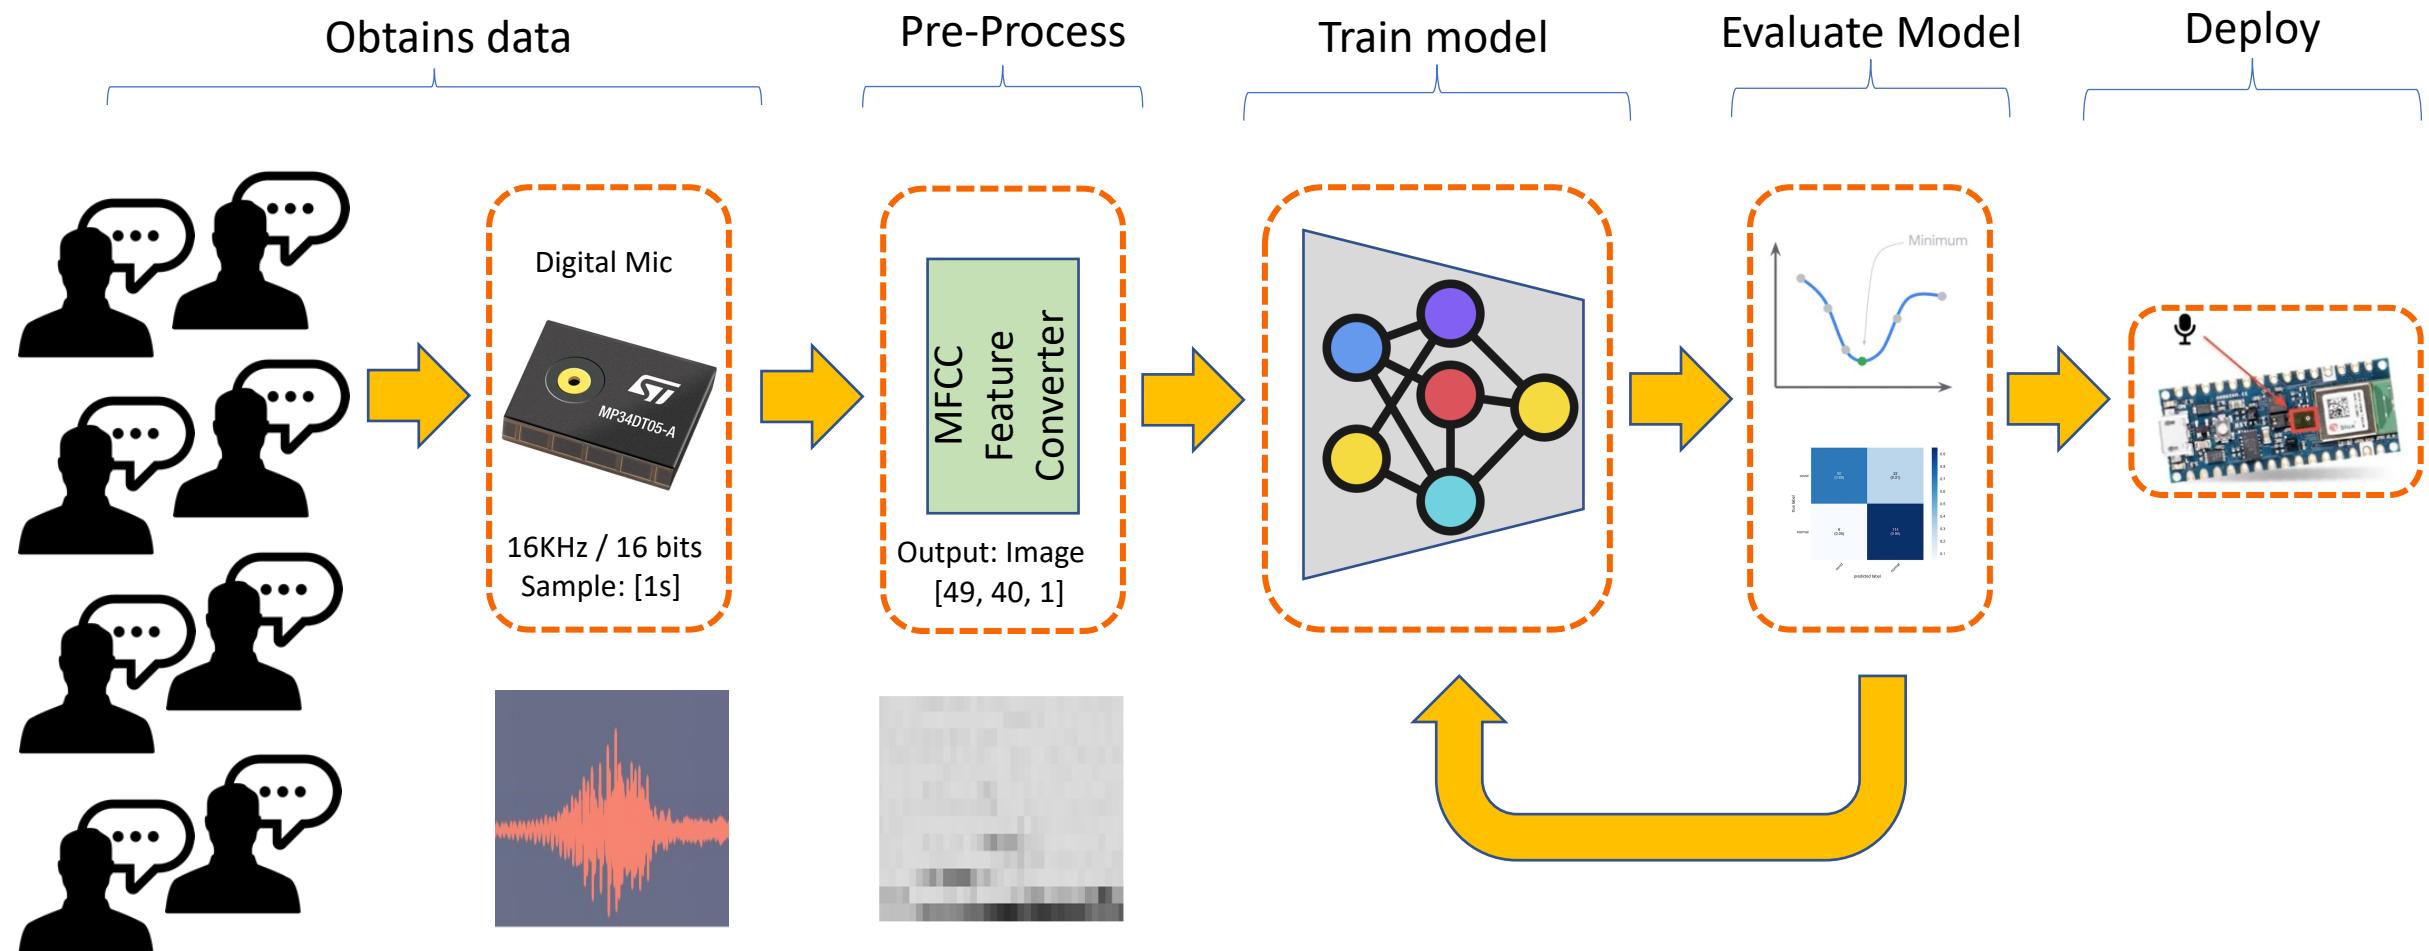
KeyWord Spotting (KWS) - Inference



KeyWord Spotting (KWS) - Model



KeyWord Spotting (KWS) – Create Model (Training)



Reading Material

Main references

- [Harvard School of Engineering and Applied Sciences - CS249r: Tiny Machine Learning](#)
- [Professional Certificate in Tiny Machine Learning \(TinyML\) – edX/Harvard](#)
- [Introduction to Embedded Machine Learning \(Coursera\)](#)
- [Text Book: "TinyML" by Pete Warden, Daniel Situnayake](#)

I want to thank [Laurence Moroney](#) from Google, Harvard professor [Vijay Janapa Reddi](#), Ph.D. student [Brian Plancher](#) and their staff for preparing the excellent material on TinyML that is the basis of this course at UNIFEI.

The IESTI01 course is part of the [TinyML4D](#), an initiative to make TinyML education available to everyone globally.

Thanks
And stay safe!

